



CS492: Probabilistic Programming Posterior Inference, Basics of Anglican, and Importance Sampling

Hongseok Yang
KAIST

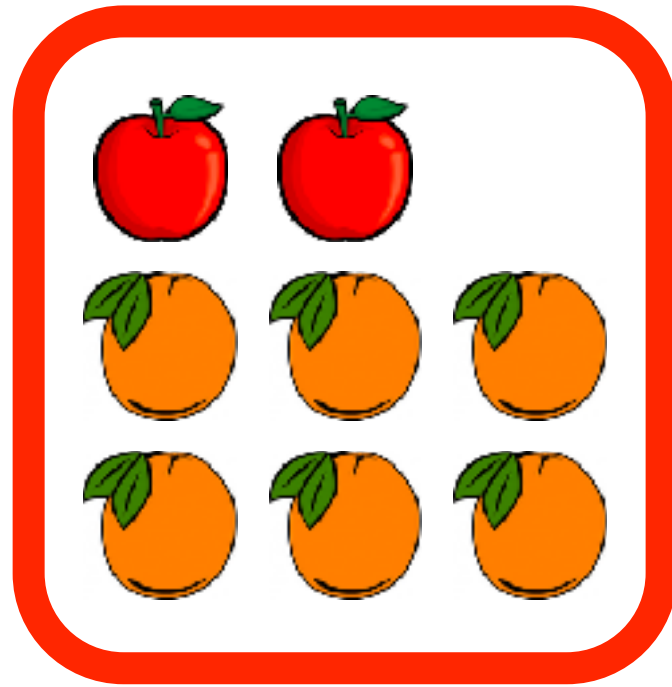
red bin



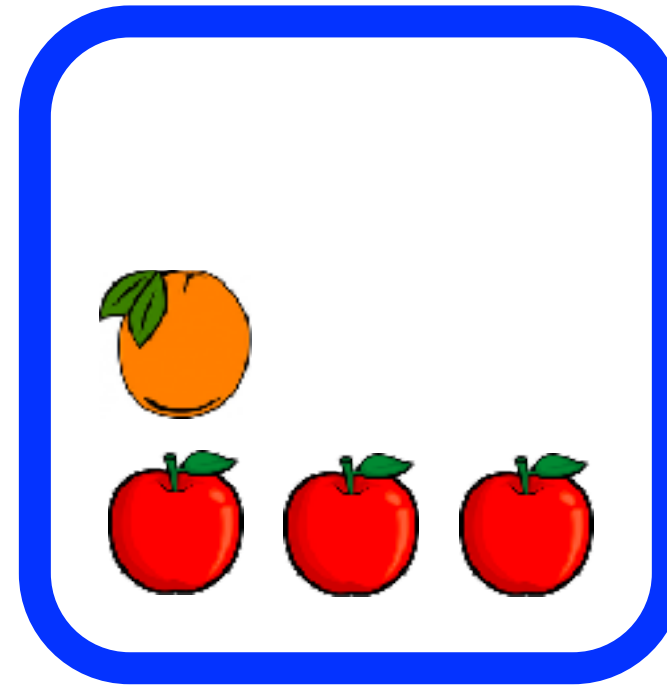
blue bin



red bin



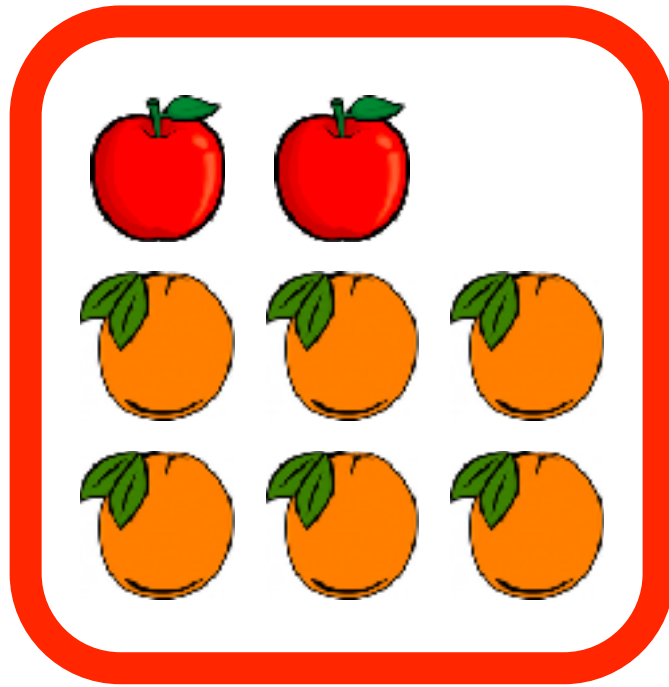
blue bin



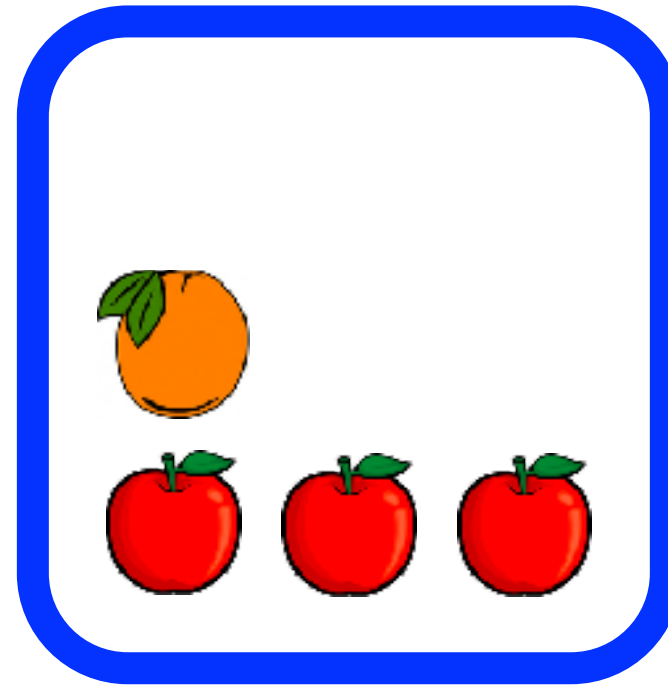
I pick a bin.

$$p(\text{red}) = 1/6 \quad p(\text{blue}) = 5/6$$

red bin



blue bin



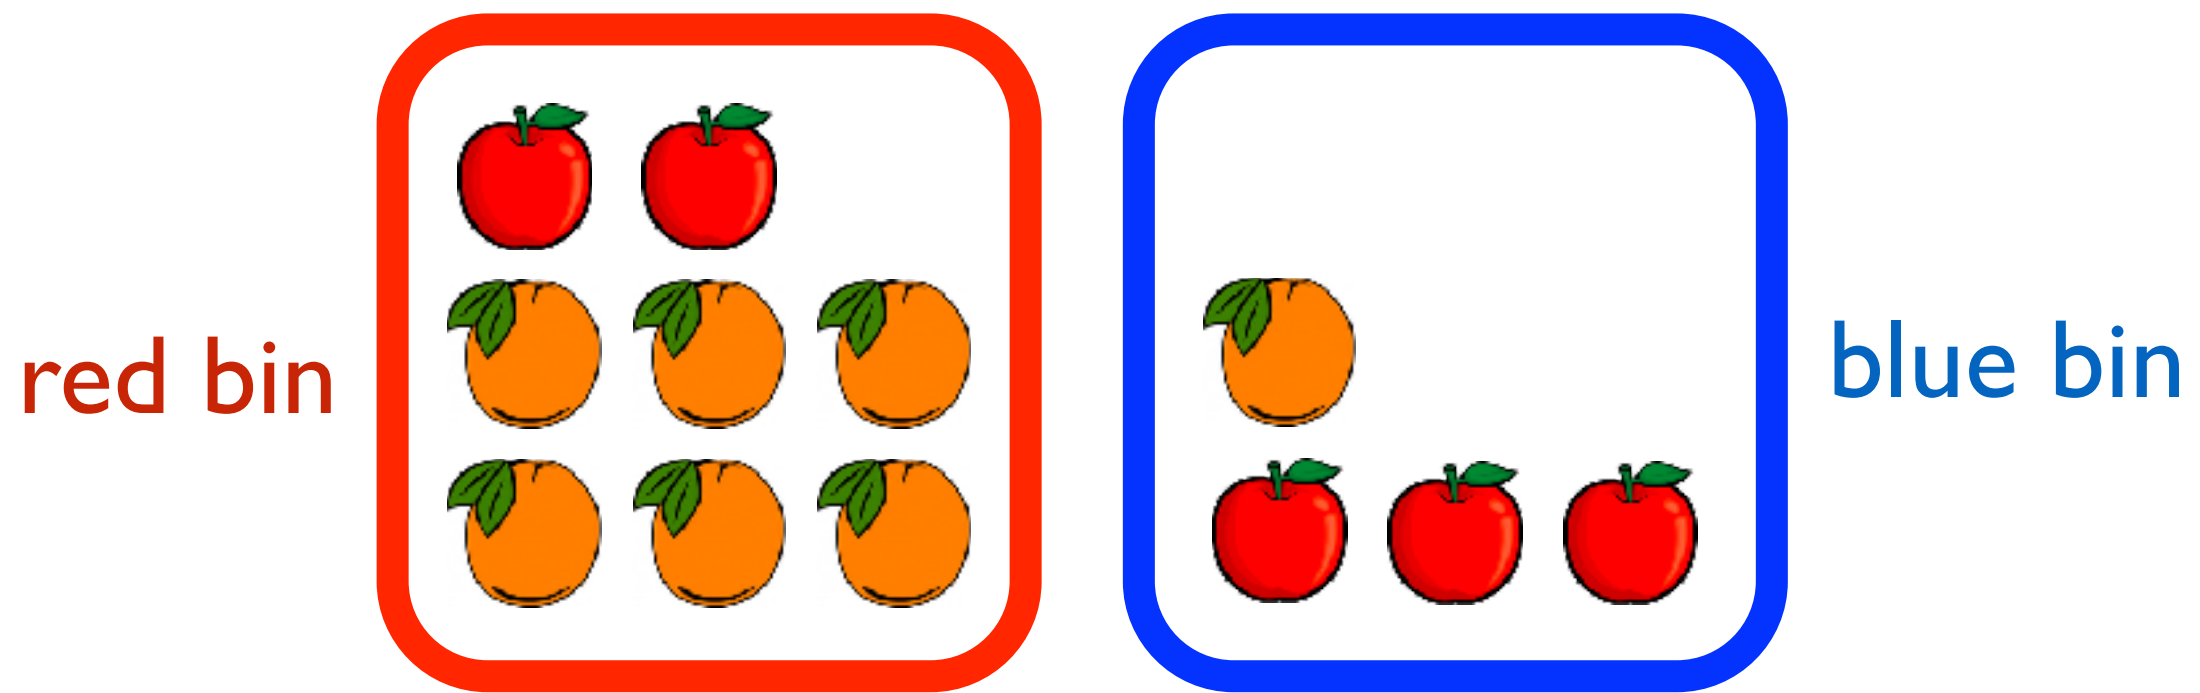
I pick a bin. Then, I choose a fruit from the bin.

$$p(\text{red}) = 1/6$$

$$p(\text{blue}) = 5/6$$

$$p(\text{apple}|\text{red}) = 2/8$$

$$p(\text{apple}|\text{blue}) = 3/4$$



I pick a bin. Then, I choose a fruit from the bin.

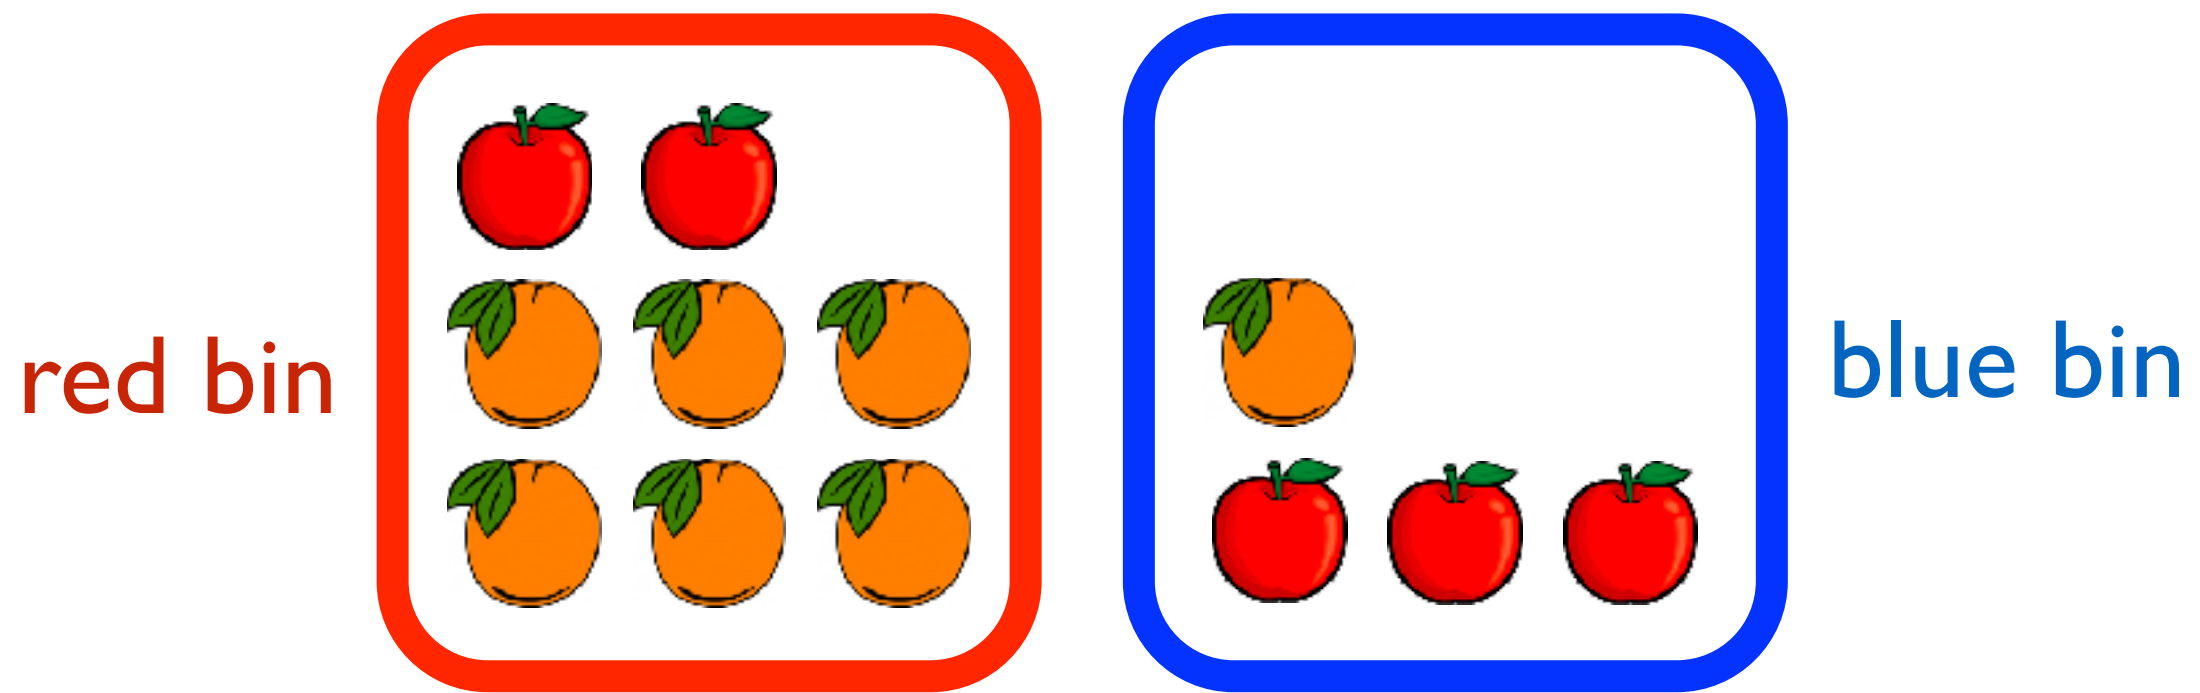
$$\begin{aligned} p(\text{red}) &= 1/6 & p(\text{blue}) &= 5/6 \\ p(\text{apple}|\text{red}) &= 2/8 & p(\text{apple}|\text{blue}) &= 3/4 \end{aligned}$$

[Q] If I pick an orange, what is the probability that I picked the blue bin?

1) $5/6$

2) $1/4$

3) $5/8$



I pick a bin. Then, I choose a fruit from the bin.

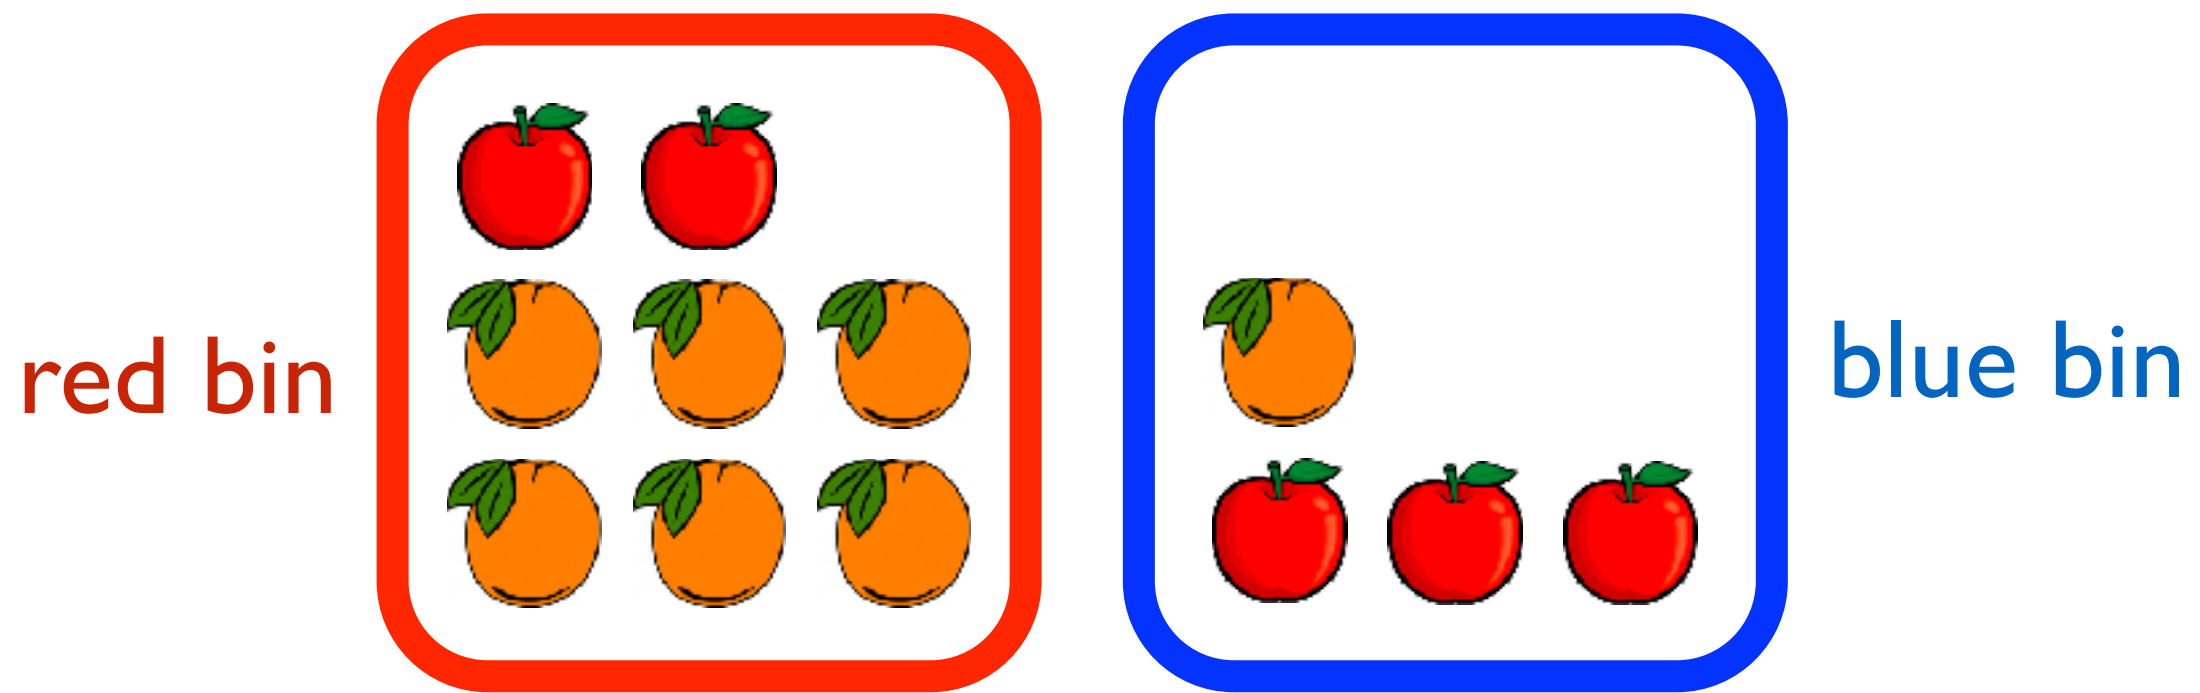
$$\begin{aligned} p(\text{red}) &= 1/6 & p(\text{blue}) &= 5/6 \\ p(\text{apple}|\text{red}) &= 2/8 & p(\text{apple}|\text{blue}) &= 3/4 \end{aligned}$$

[Q] If I pick an orange, what is the probability that I picked the blue bin?

1) $5/6$

2) $1/4$

3) $5/8$



I pick a bin. Then, I choose a fruit from the bin.


$$\begin{aligned} p(\text{red}) &= 1/6 & p(\text{blue}) &= 5/6 \\ p(\text{apple}|\text{red}) &= 2/8 & p(\text{apple}|\text{blue}) &= 3/4 \end{aligned}$$

[Q] If I pick an orange, what is the probability that I picked the blue bin?

1) $5/6$

2) $1/4$

3) $5/8$



I pick a bin. Then, I choose a fruit from the bin.

$$p(\text{red}) = 1/6$$

$$p(\text{blue}) = 5/6$$

$$p(\text{apple}|\text{red}) = 2/8$$

$$p(\text{apple}|\text{blue}) = 3/4$$

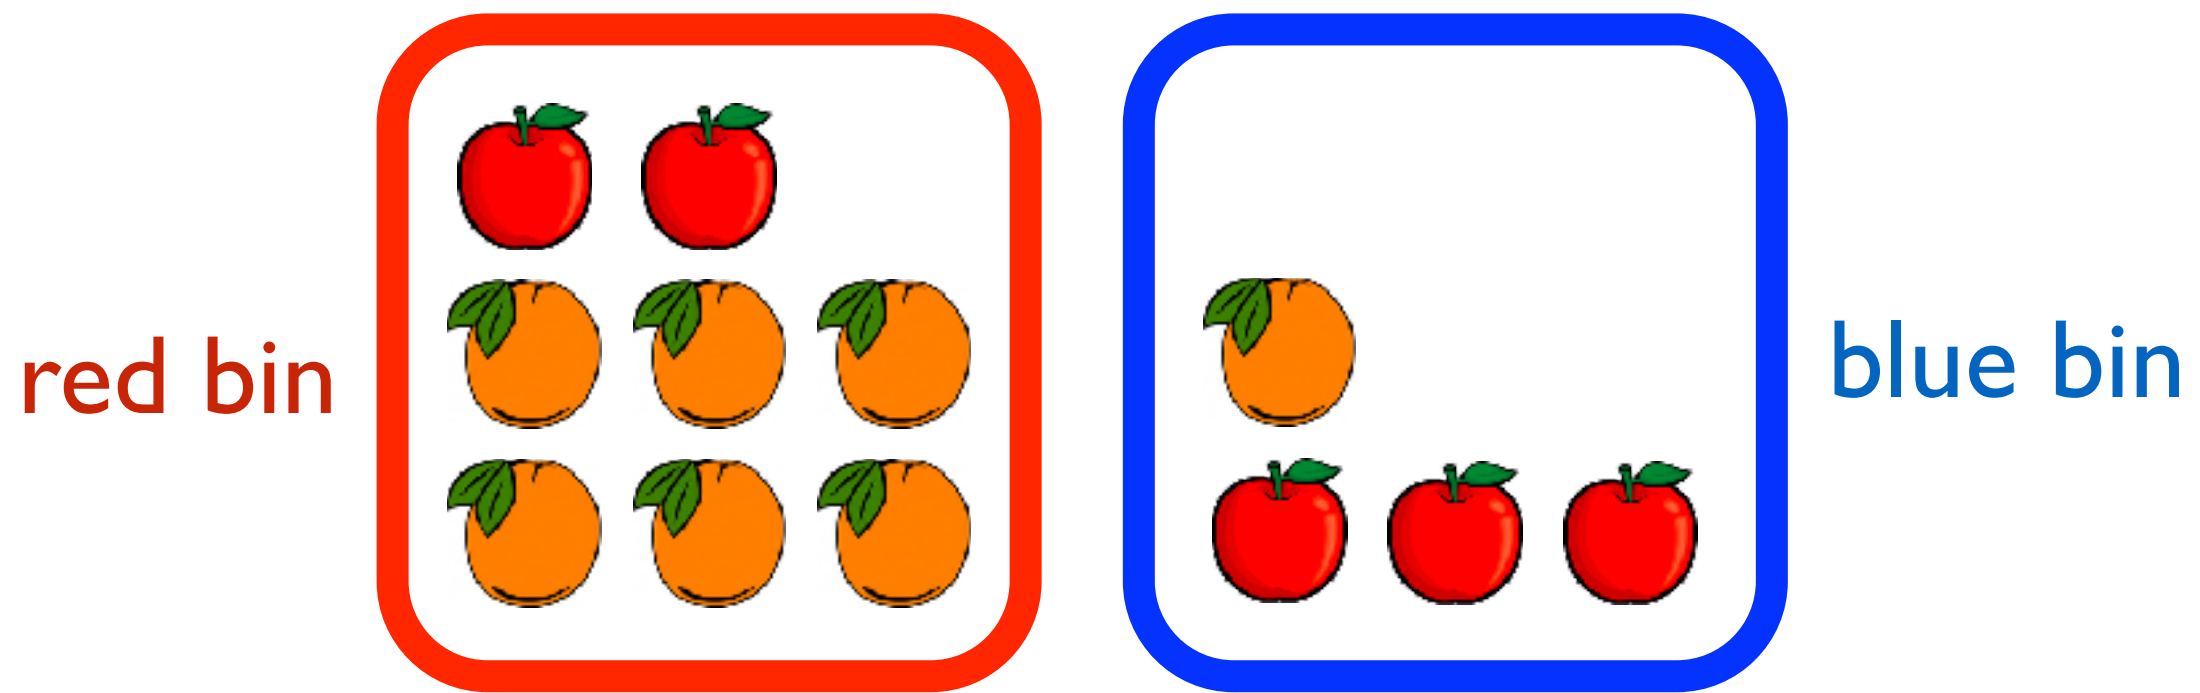
[Q] $p(\text{orange}|\text{red}) = 3/4$ $p(\text{orange}|\text{blue}) = 1/4$

that I picked the blue bin?

1) $5/6$

2) $1/4$

3) $5/8$



I pick a bin. Then, I choose a fruit from the bin.

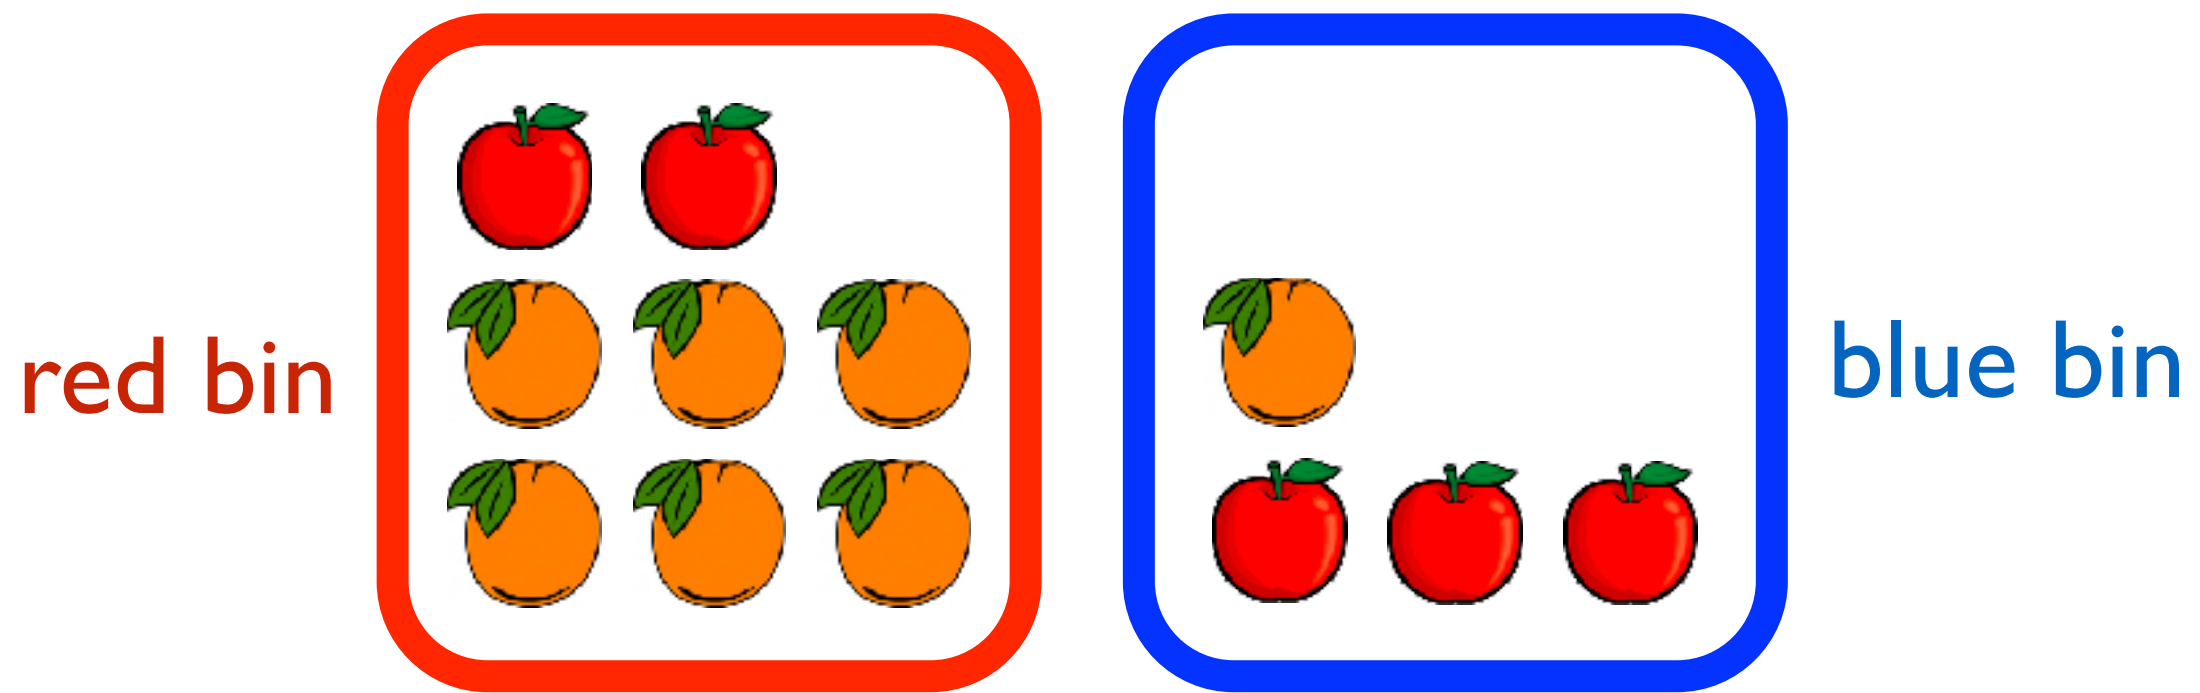
$$\begin{aligned} p(\text{red}) &= 1/6 & p(\text{blue}) &= 5/6 \\ p(\text{apple}|\text{red}) &= 2/8 & p(\text{apple}|\text{blue}) &= 3/4 \end{aligned}$$

[Q] If I pick an orange, what is the probability that I picked the blue bin?

1) 5/6

2) 1/4

3) 5/8



I pick a bin. Then, I choose a fruit from the bin.

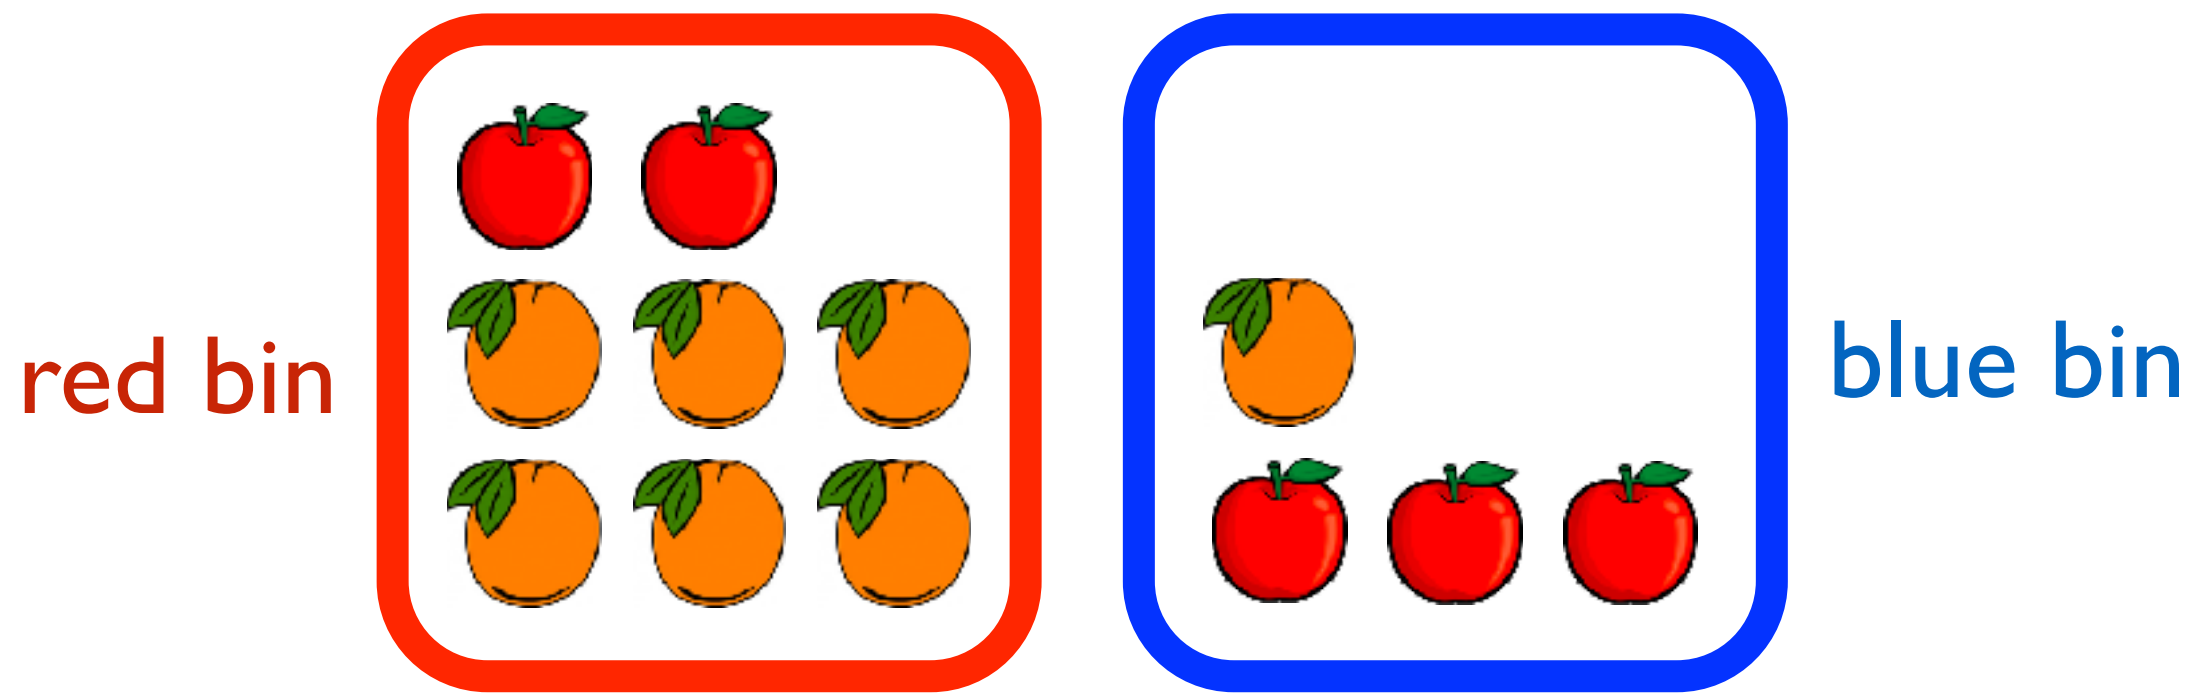
$$\begin{aligned} p(\text{red}) &= 1/6 & p(\text{blue}) &= 5/6 \\ p(\text{apple}|\text{red}) &= 2/8 & p(\text{apple}|\text{blue}) &= 3/4 \end{aligned}$$

[Q] If I pick an orange, what is the probability that I picked the blue bin?

1) $5/6$

2) $1/4$

3) $5/8$



I pick a bin. Then, I choose a fruit from the bin.

$$\begin{aligned} p(\text{red}) &= 1/6 & p(\text{blue}) &= 5/6 \\ p(\text{apple}|\text{red}) &= 2/8 & p(\text{apple}|\text{blue}) &= 3/4 \end{aligned}$$

[Q] If I pick an orange, what is the probability that I picked the blue bin?

1) $5/6$

2) $1/4$

3) $5/8$

Learning outcome

- Can describe prior, likelihood, posterior, Bayes' rule.
- Can solve the puzzle using Bayes' rule
- Can express/solve the puzzle in Anglican.
- Can explain importance sampling.

We will use discrete probabilities mostly.

Review of discrete probability, and posterior inference

- Consider random variables x, y, z, \dots having values in countable sets, such as $\{\text{true}, \text{false}\}$ and \mathbb{N} .

- Consider random variables x, y, z, \dots having values in countable sets, such as $\{\text{true}, \text{false}\}$ and \mathbb{N} .
- Probability p assigns numbers between 0 and 1 for all possible value assignments of some variables.

$$\begin{array}{ll} p(x=0, y=0) = 1/24 & p(x=0, y=1) = 3/24 \\ p(x=1, y=0) = 5/24 & p(x=1, y=1) = 15/24 \\ p(x=0) = 4/24 & p(x=1) = 20/24 \end{array}$$

- Consider random variables x, y, z, \dots having values in countable sets, such as $\{\text{true}, \text{false}\}$ and \mathbb{N} .
- Probability p assigns numbers between 0 and 1 for all possible value assignments of some variables.

$$\begin{array}{ll}
 p(x=0, y=0) = 1/24 & p(x=0, y=1) = 3/24 \\
 p(x=1, y=0) = 5/24 & p(x=1, y=1) = 15/24 \\
 p(x=0) = 4/24 & p(x=1) = 20/24
 \end{array}$$

[Requirement 1] $\sum_{v,w} p(x=v, y=w) = 1$.

- Consider random variables x, y, z, \dots having values in countable sets, such as $\{\text{true}, \text{false}\}$ and \mathbb{N} .
- Probability p assigns numbers between 0 and 1 for all possible value assignments of some variables.

$$\begin{array}{ll}
 p(x=0, y=0) = 1/24 & p(x=0, y=1) = 3/24 \\
 p(x=1, y=0) = 5/24 & p(x=1, y=1) = 15/24 \\
 p(x=0) = 4/24 & p(x=1) = 20/24
 \end{array}$$

[Requirement 1] $\sum_{v,w} p(x=v, y=w) = 1.$

- Consider random variables x, y, z, \dots having values in countable sets, such as $\{\text{true}, \text{false}\}$ and \mathbb{N} .
- Probability p assigns numbers between 0 and 1 for all possible value assignments of some variables.

$$\begin{array}{ll}
 p(x=0, y=0) = 1/24 & p(x=0, y=1) = 3/24 \\
 p(x=1, y=0) = 5/24 & p(x=1, y=1) = 15/24 \\
 p(x=0) = 4/24 & p(x=1) = 20/24
 \end{array}$$

[Requirement 1] $\sum_{v,w} p(x=v, y=w) = 1.$

- Consider random variables x, y, z, \dots having values in countable sets, such as $\{\text{true}, \text{false}\}$ and \mathbb{N} .
- Probability p assigns numbers between 0 and 1 for all possible value assignments of some variables.

$$\begin{array}{ll}
 p(x=0, y=0) = 1/24 & p(x=0, y=1) = 3/24 \\
 p(x=1, y=0) = 5/24 & p(x=1, y=1) = 15/24 \\
 p(x=0) = 4/24 & p(x=1) = 20/24
 \end{array}$$

[Requirement 1] $\sum_{v,w} p(x=v, y=w) = 1$.

[Requirement 2] $p(x=v) = \sum_w p(x=v, y=w)$.

- Consider random variables x, y, z, \dots having values in countable sets, such as $\{\text{true}, \text{false}\}$ and \mathbb{N} .
- Probability p assigns numbers between 0 and 1 for all possible value assignments of some variables.

$$\begin{array}{ll}
 p(x=0, y=0) = 1/24 & p(x=0, y=1) = 3/24 \\
 p(x=1, y=0) = 5/24 & p(x=1, y=1) = 15/24 \\
 p(x=0) = 4/24 & p(x=1) = 20/24
 \end{array}$$

[Requirement 1] $\sum_{v,w} p(x=v, y=w) = 1$.

[Requirement 2] $p(x=v) = \sum_w p(x=v, y=w)$.

- Consider random variables x, y, z, \dots having values in countable sets, such as $\{\text{true}, \text{false}\}$ and \mathbb{N} .
- Probability p assigns numbers between 0 and 1 for all possible value assignments of some variables.

$$\begin{array}{ll}
 p(x=0, y=0) = 1/24 & p(x=0, y=1) = 3/24 \\
 p(x=1, y=0) = 5/24 & p(x=1, y=1) = 15/24 \\
 p(x=0) = 4/24 & p(x=1) = 20/24
 \end{array}$$

[Requirement 1] $\sum_{v,w} p(x=v, y=w) = 1$.

[Requirement 2] $p(x=v) = \sum_w p(x=v, y=w)$.

[Q] Compute $p(y=0)$ and $p(y=1)$.

- Consider random variables x, y, z, \dots having values in countable sets, such as $\{\text{true}, \text{false}\}$ and \mathbb{N} .
- Probability p assigns numbers between 0 and 1 to all possible value assignments of so $p(x=v), p(y=w)$.

Enough.

Determines

$p(x=v), p(y=w)$.

$$\begin{array}{ll}
 p(x=0, y=0) = 1/24 & p(x=0, y=1) = 3/24 \\
 p(x=1, y=0) = 5/24 & p(x=1, y=1) = 15/24 \\
 p(x=0) = 4/24 & p(x=1) = 20/24
 \end{array}$$

[Requirement 1] $\sum_{v,w} p(x=v, y=w) = 1$.

[Requirement 2] $p(x=v) = \sum_w p(x=v, y=w)$.

[Q] Compute $p(y=0)$ and $p(y=1)$.

Conditional probability

$$p(x=v \mid y=w) =_{\text{def}} \frac{p(x=v, y=w)}{p(y=w)}$$

Says the prob. of $x=v$ conditioned on $y=w$.

Conditional probability

$$p(x=v \mid y=w) =_{\text{def}} \frac{p(x=v, y=w)}{p(y=w)}$$

Says the prob. of $x=v$ conditioned on $y=w$.

[Lemma 1] $\sum_v p(x=v \mid y=w) = 1$.

Conditional probability

$$p(x=v \mid y=w) =_{\text{def}} \frac{p(x=v, y=w)}{p(y=w)}$$

Says the prob. of $x=v$ conditioned on $y=w$.

[Lemma 1] $\sum_v p(x=v \mid y=w) = 1$.

[Lemma 2] (Bayes' rule)

$$p(x=v \mid y=w) = \frac{p(y=w \mid x=v) \times p(x=v)}{p(y=w)}$$

Conditional probability

$$p(x=v \mid y=w) =_{\text{def}} \frac{p(x=v, y=w)}{p(y=w)}$$

Says the prob. of $x=v$ conditioned on $y=w$.

[Lemma 1] $\sum_v p(x=v \mid y=w) = 1$.

[Lemma 2] (Bayes' rule)

$$p(x \mid y) = \frac{p(y \mid x) \times p(x)}{p(y)}$$

In sloppy but simpler popular notation.

Conditional probability

$$p(x=v \mid y=w) =_{\text{def}} \frac{p(x=v, y=w)}{p(y=w)}$$

Says the prob. of $x=v$ conditioned on $y=w$.

[Lemma 1] $\sum_v p(x=v \mid y=w) = 1$.

[Lemma 2] (Bayes' rule)

$$p(x \mid y) = \frac{p(y \mid x) \times p(x)}{p(y)}$$

In sloppy but simpler popular notation.

Conditional probability

$$p(x=v \mid y=w) =_{\text{def}} \frac{p(x=v, y=w)}{p(y=w)}$$

Says the prob. of $x=v$ conditioned on $y=w$.

[Lemma 1] $\sum_v p(x=v \mid y=w) = 1$.

[Q] Prove both lemmas.

[Lemma 2] (Bayes' rule)

$$p(x \mid y) = \frac{p(y \mid x) \times p(x)}{p(y)}$$

In sloppy but simpler popular notation.

Bayes' rule

$$p(x | y) = \frac{p(y | x) \times p(x)}{p(y)}$$

- Trivial fact. But super famous. Why?

Bayes' rule

$$p(x | y) = \frac{p(y | x) \times p(x)}{p(y)}$$

- Trivial fact. But super famous. Why?
- Says how to combine prior knowledge with observed data consistently.

y — observation

Bayes' rule

$$p(x | y) = \frac{p(y | x) \times p(x)}{p(y)}$$

- Trivial fact. But super famous. Why?
- Says how to combine prior knowledge with observed data consistently.

y — observation
 x — target latent
variable

Bayes' rule

$$p(x | y) = \frac{p(y | x) \times p(x)}{p(y)}$$

- Trivial fact. But super famous. Why?
- Says how to combine prior knowledge with observed data consistently.

y — observation
 x — target latent
variable

Bayes' rule

$$p(x | y) = \frac{p(y | x) \times p(x)}{p(y)}$$

- Trivial fact. But super famous. Why?
- Says how to combine prior knowledge with observed data consistently.
- Typically, $p(x)$ & $p(y|x)$ specified (not $p(x,y)$).

y — observation
 x — target latent
variable

Bayes' rule

$$p(x | y) = \frac{p(y | x) \times p(x)}{p(y)}$$

prior
distribution



- Trivial fact. But super famous. Why?
- Says how to combine **prior knowledge** with observed data consistently.
- Typically, $p(x)$ & $p(y|x)$ specified (not $p(x,y)$).

y — observation
 x — target latent
variable

Bayes' rule

likelihood

prior
distribution

$$p(x | y) = \frac{p(y | x) \times p(x)}{p(y)}$$

- Trivial fact. But super famous. Why?
- Says how to combine prior knowledge with **observed data** consistently.
- Typically, $p(x)$ & $p(y|x)$ specified (not $p(x,y)$).

y — observation
 x — target latent
variable

Bayes' rule

likelihood

prior
distribution

$$p(x | y) = \frac{p(y | x) \times p(x)}{p(y)}$$

posterior
distribution

- Trivial fact. But super famous. Why?
- Says how to combine prior knowledge with observed data consistently.
- Typically, $p(x)$ & $p(y|x)$ specified (not $p(x,y)$).

y — observation
 x — target latent
variable

Bayes' rule

$$p(x | y) = \frac{p(y | x) \times p(x)}{p(y)}$$

posterior distribution

likelihood

prior distribution

marginal likelihood

- Trivial fact. But super famous. Why?
- Says how to combine prior knowledge with observed data consistently.
- Typically, $p(x)$ & $p(y|x)$ specified (not $p(x,y)$).

y — observation
 x — target latent
variable

Bayes' rule

$$p(x | y) = \frac{p(y | x) \times p(x)}{p(y)}$$

posterior distribution

likelihood

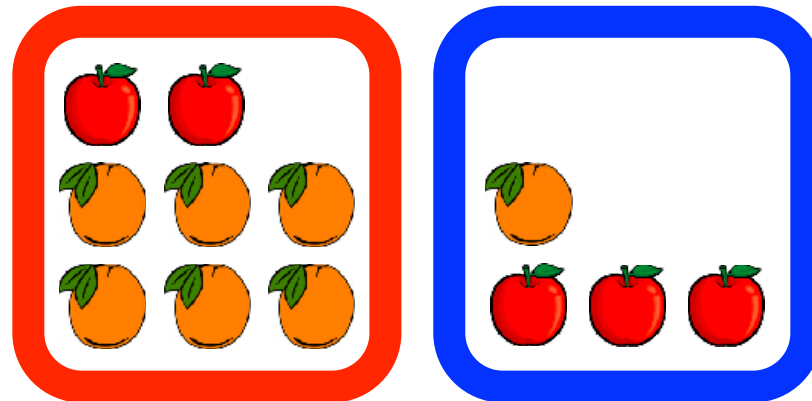
prior distribution

marginal likelihood

posterior \propto likelihood \times prior

- Trivial fact. But super famous. Why?
- Says how to combine prior knowledge with observed data consistently.
- Typically, $p(x)$ & $p(y|x)$ specified (not $p(x,y)$).

Puzzle again



I pick a bin. Then, I choose a fruit from the bin.

$$\begin{aligned} p(\text{red}) &= 1/6 & p(\text{blue}) &= 5/6 \\ p(\text{apple}|\text{red}) &= 2/8 & p(\text{apple}|\text{blue}) &= 3/4 \end{aligned}$$

[Q] If I pick an orange, what is the probability that I picked the blue bin?

Exercise

A bag contains one ball, either white with prob $1/5$ or black with prob $4/5$. An additional white ball is put in, and the bag is shaken. Then, a ball is drawn, which proves to be white. What is now the chance of drawing a white ball?

Modified Ex 3.12 from MacKay's Info.Th. book

Posterior inference

- Computation of $p(x|y)$ given i) $p(y|x)$ and $p(x)$ and ii) an observed value w of y .
- Bayes' rule and Req 2 give an algorithm:

$$\begin{aligned} p(x \mid y=w) &= \frac{p(y=w \mid x) \times p(x)}{p(y=w)} \\ &= \frac{p(y=w \mid x) \times p(x)}{\sum_v p(x=v, y=w)} \end{aligned}$$

Posterior inference

- Computation of $p(x|y)$ given i) $p(y|x)$ and $p(x)$ and ii) an observed value w of y .
- Bayes' rule and Req 2 give an algorithm:

$$\begin{aligned} p(x \mid y=w) &= \frac{p(y=w \mid x) \times p(x)}{p(y=w)} \\ &= \frac{p(y=w \mid x) \times p(x)}{\sum_v p(x=v, y=w)} \end{aligned}$$

Big sum for realistic models. Inefficient.

Approximate posterior inference

- Approximates posterior $p(x|y)$ using a set of samples or a simpler distribution.
- Commonly used in practice.
- Anglican implements many such algorithms.

Expectation

$$\mathbb{E}_{p(x)}[f(x)] = \sum_x p(x)f(x)$$

1. Linearity:

$$\mathbb{E}_{p(x)}[\alpha f(x) + \beta g(x)] = \alpha \mathbb{E}_{p(x)}[f(x)] + \beta \mathbb{E}_{p(x)}[g(x)]$$

2. Independent random variables:

$$\mathbb{E}_{p(x)p(y)}[f(x)g(y)] = \mathbb{E}_{p(x)}[f(x)]\mathbb{E}_{p(y)}[g(y)]$$

Expectation

$$\mathbb{E}_{p(x)}[f(x)] = \sum_x p(x)f(x)$$

1. Linearity:

$$\mathbb{E}_{p(x)}[\alpha f(x) + \beta g(x)] = \alpha \mathbb{E}_{p(x)}[f(x)] + \beta \mathbb{E}_{p(x)}[g(x)]$$

2. Independent random variables:

$$\mathbb{E}_{p(x)p(y)}[f(x)g(y)] = \mathbb{E}_{p(x)}[f(x)]\mathbb{E}_{p(y)}[g(y)]$$

[Q] A coin with probability p of coming up heads. N coin tosses. Mean of the number of heads? Variance?

Conditioning and posterior inference in Anglican

Conditioning in Anglican

In Anglican, we condition a model by observed random variables using the observe construct:

(observe distribution-object observed-value)

Examples:

```
(observe (flip p) true)
```

```
(observe  
  (categorical  
    {:blue p, :red q, :green r})  
  :blue)
```

[Q] Write an Anglican query for our puzzle using categorical distribution.

```
(defquery puz1 [fruit]
```

```
  (let [bin
```



```
  ]
```



```
  bin))
```


[Q] Write an Anglican query for our puzzle using categorical distribution.

```
(defquery puz1 [fruit]
  (let [bin
```



```
]
```



```
bin))
```

```
(observe
  (categorical
    {:blue p, :red q, :green r})
  :blue)
```

[Q] Write an Anglican query for our puzzle using categorical distribution.

```
(defquery puz1 [fruit]
  (let [bin (sample (categorical
                    { :red (/ 1 6),
                      :blue (/ 5 6) }))]
    (if (= bin :red)
```



```
bin))
```

```
(observe
  (categorical
    { :blue p, :red q, :green r })
  :blue)
```

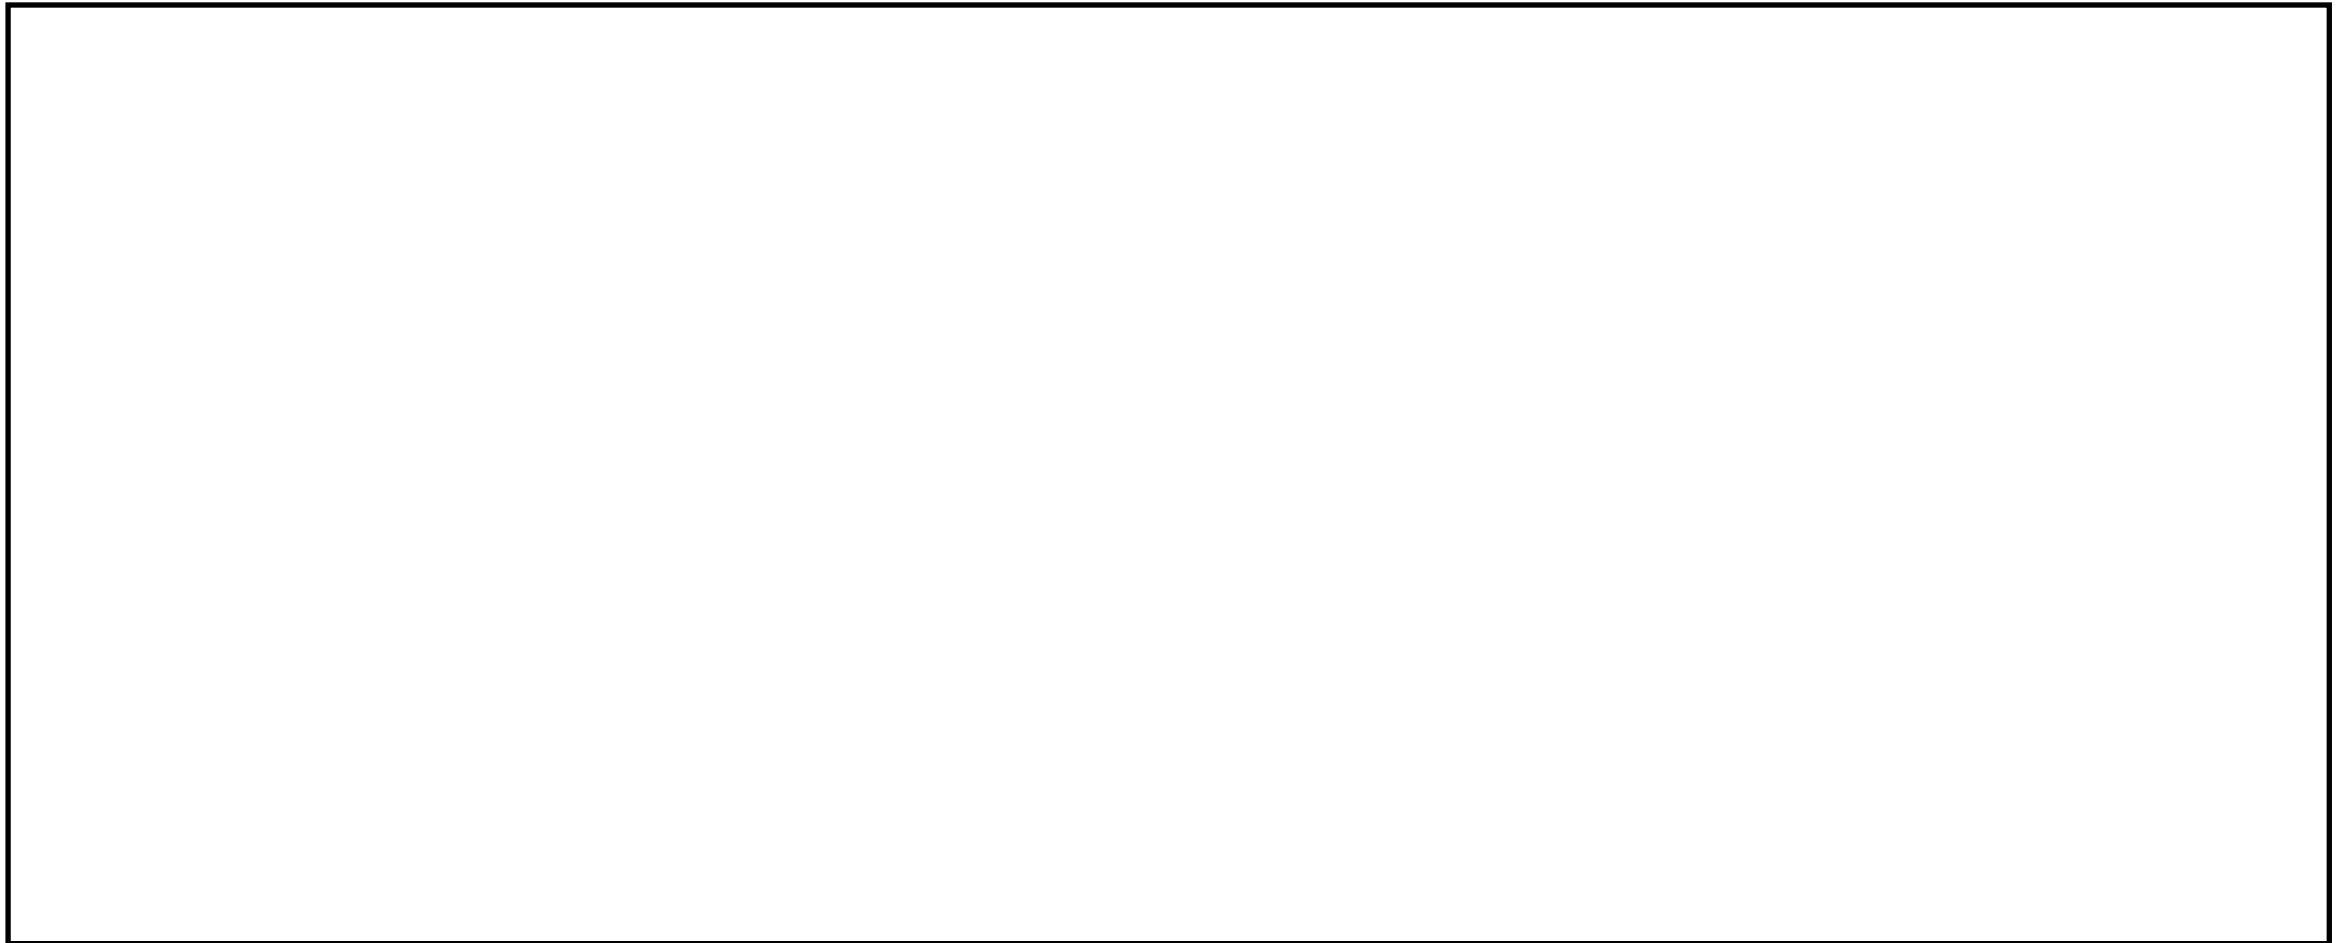
[Q] Write an Anglican query for our puzzle using categorical distribution.

```
(defquery puz1 [fruit]
  (let [bin (sample (categorical
                    { :red (/ 1 6),
                      :blue (/ 5 6) }) )]

    (if (= bin :red)
        (observe (categorical
                  { :apple (/ 2 8),
                    :orange (/ 6 8) })
                fruit)
        (observe (categorical
                  { :apple (/ 3 4),
                    :orange (/ 1 4) })
                fruit))

    bin))
```

We perform approximate posterior inference using the importance-sampling algo. of Anglican.



We perform approximate posterior inference using the importance-sampling algo. of Anglican.

```
(def x (doquery :importance puz1 [:orange]))
```

We perform approximate posterior inference using the importance-sampling algo. of Anglican.

```
(def x (doquery :importance puz1 [:orange]))
```



Anglican function.

Performs inference.

Returns a lazy infinite list of Clojure maps.

We perform approximate posterior inference using the importance-sampling algo. of Anglican.

```
(def x (doquery :importance puz1 [:orange]))  
(println (first x))
```

We perform approximate posterior inference using the importance-sampling algo. of Anglican.

```
(def x (doquery :importance puz1 [:orange]))  
(println (first x))
```

```
{:log-weight -1.3862943611198906,  
 :result :blue, :predicts []}
```


We perform approximate posterior inference using the importance-sampling algo. of Anglican.

```
(def x (doquery :importance puz1 [:orange]))
```

We perform approximate posterior inference using the importance-sampling algo. of Anglican.

```
(def x (doquery :importance puz1 [:orange]))  
(def y (take 10000 x))
```

We perform approximate posterior inference using the importance-sampling algo. of Anglican.

```
(def x (doquery :importance puz1 [:orange]))  
(def y (take 10000 x))  
(println (count y))  
(println (first (rest y)))
```

We perform approximate posterior inference using the importance-sampling algo. of Anglican.

```
(def x (doquery :importance puz1 [:orange]))  
(def y (take 10000 x))  
(println (count y))  
(println (first (rest y)))
```

```
10000  
{:log-weight -1.3862943611198906,  
 :result :blue, :predicts []}
```

We perform approximate posterior inference using the importance-sampling algo. of Anglican.

```
(def x (doquery :importance puz1 [:orange]))  
(def y (take 10000 x))
```

We perform approximate posterior inference using the importance-sampling algo. of Anglican.

```
(def x (doquery :importance puz1 [:orange]))  
(def y (take 10000 x))  
  
(defn f [m] (exp (:log-weight m)))  
(defn g [m]  
  (if (= (:result m) :blue) (f m) 0.0))
```

We perform approximate posterior inference using the importance-sampling algo. of Anglican.

```
(def x (doquery :importance puz1 [:orange]))
(def y (take 10000 x))

(defn f [m] (exp (:log-weight m)))
(defn g [m]
  (if (= (:result m) :blue) (f m) 0.0))

(/ (reduce + 0.0 (map g y))
   (reduce + 0.0 (map f y)))
```

We perform approximate posterior inference using the importance-sampling algo. of Anglican.

```
(def x (doquery :importance puz1 [:orange]))
(def y (take 10000 x))

(defn f [m] (exp (:log-weight m)))
(defn g [m]
  (if (= (:result m) :blue) (f m) 0.0))

(/ (reduce + 0.0 (map g y))
   (reduce + 0.0 (map f y)))
```

[Q] Does anyone see what goes on here?

We perform approximate posterior inference using the importance-sampling algo. of Anglican.

```
(def x (doquery :importance puz1 [:orange]))
(def y (take 10000 x))

(defn f [m] (exp (:log-weight m)))
(defn g [m]
  (if (= (:result m) :blue) (f m) 0.0))

(/ (reduce + 0.0 (map g y))
   (reduce + 0.0 (map f y)))
```

[Q] Does anyone see what goes on here?

[A] Portion of (weighted) blue samples among all (weighted) samples.

Likelihood weighted importance sampling

[Goal] Estimate $\mathbb{E}_{p(\mathbf{x}|y)}[f(\mathbf{x})]$ for a given f .

Likelihood weighted importance sampling

[Goal] Estimate $\mathbb{E}_{p(x|y)}[f(x)]$ for a given f .


posterior

Likelihood weighted importance sampling

[Goal] Estimate $\mathbb{E}_{p(x|y)}[f(x)]$ for a given f .

I. Sample x_1, \dots, x_N from **prior** $p(x)$.

Likelihood weighted importance sampling

[Goal] Estimate $\mathbb{E}_{p(x|y)}[f(x)]$ for a given f .

1. Sample x_1, \dots, x_N from prior $p(x)$.
2. Compute **weight** $w_i = p(y|x_i)$ for each i .

Likelihood weighted importance sampling

[Goal] Estimate $\mathbb{E}_{p(x|y)}[f(x)]$ for a given f .

1. Sample x_1, \dots, x_N from prior $p(x)$.
2. Compute weight $w_i = p(y|x_i)$ for each i .
3. Return **weighted** avg. $(\sum_i w_i x f(x_i)) / \sum_j w_j$.

Likelihood weighted importance sampling

[Goal] Estimate $\mathbb{E}_{p(x|y)}[f(x)]$ for a given f .

1. Sample x_1, \dots, x_N from prior $p(x)$.
2. Compute weight $w_i = p(y|x_i)$ for each i .
3. Return weighted avg. $(\sum_i w_i x f(x_i)) / \sum_j w_j$.

[puzl] $f(x)=0$ if $(:result\ x)$ is `:red`. If `:blue`, $f(x)=1$.

Likelihood weighted importance sampling

[Goal] Estimate $\mathbb{E}_{p(x|y)}[f(x)]$ for a given f .

1. Sample x_1, \dots, x_N from prior $p(x)$.
2. Compute weight $w_i = p(y|x_i)$ for each i .
3. Return weighted avg. $(\sum_i w_i x f(x_i)) / \sum_j w_j$.

[puzl] $f(x)=0$ if $(\text{:result } x)$ is :red . If :blue , $f(x)=1$.

[Q1] Why is this a sensible algorithm?

Likelihood weighted importance sampling

[Goal] Estimate $\mathbb{E}_{p(x|y)}[f(x)]$ for a given f .

1. Sample x_1, \dots, x_N from prior $p(x)$.
2. Compute weight $w_i = p(y|x_i)$ for each i .
3. Return weighted avg. $(\sum_i w_i x f(x_i)) / \sum_j w_j$.

[puzl] $f(x)=0$ if $(\text{:result } x)$ is :red . If :blue , $f(x)=1$.

[Q1] Why is this a sensible algorithm?

[Q2] How to implement 1 & 2 for Anglican queries?

[Input] N and an Anglican query Q .

[Output] Weighted samples $(w_1, s_1), \dots, (w_N, s_N)$.

[Input] N and an Anglican query Q .

[Output] Weighted samples $(w_1, s_1), \dots, (w_N, s_N)$.

Run Q as follows for N times.

[Input] N and an Anglican query Q .

[Output] Weighted samples $(w_1, s_1), \dots, (w_N, s_N)$.

Run Q as follows for N times.

1. Use a global variable w initialised to 1.

[Input] N and an Anglican query Q .

[Output] Weighted samples $(w_1, s_1), \dots, (w_N, s_N)$.

Run Q as follows for N times.

1. Use a global variable w initialised to 1.
2. Run Q as usual except for sample & observe.

[Input] N and an Anglican query Q .

[Output] Weighted samples $(w_1, s_1), \dots, (w_N, s_N)$.

Run Q as follows for N times.

1. Use a global variable w initialised to 1.
2. Run Q as usual except for `sample` & observe.
 - a. (`sample dist`): draw a sample from *dist*.

[Input] N and an Anglican query Q .

[Output] Weighted samples $(w_1, s_1), \dots, (w_N, s_N)$.

Run Q as follows for N times.

1. Use a global variable w initialised to 1.
2. Run Q as usual except for `sample` & `observe`.
 - a. (`sample dist`): draw a sample from *dist*.
 - b. (`observe dist v`): update $w := w \times p_{dist}(v)$.

[Input] N and an Anglican query Q .

[Output] Weighted samples $(w_1, s_1), \dots, (w_N, s_N)$.

Run Q as follows for N times.

1. Use a global variable w initialised to 1.
2. Run Q as usual except for `sample` & `observe`.
 - a. (`sample dist`): draw a sample from *dist*.
 - b. (`observe dist v`): update $w := w \times p_{dist}(v)$.
3. Return w and the result s of Q .

fruit = :orange

```
(defquery puz1 [fruit]
  (let [bin (sample
              (categorical
               {:red (/ 1 6),
                :blue (/ 5 6)})])
        (if (= bin :red)
            (observe (categorical
                     {:apple (/ 2 8),
                      :orange (/ 6 8)})
                    fruit)
            (observe (categorical
                     {:apple (/ 3 4),
                      :orange (/ 1 4)})
                    fruit))
        bin))
```

fruit = :orange

```
(defquery puz1 [fruit]
  (let [bin (sample
              (categorical
               {:red (/ 1 6),
                :blue (/ 5 6)})])
        (if (= bin :red)
            (observe (categorical
                     {:apple (/ 2 8),
                      :orange (/ 6 8)})
                    fruit)
            (observe (categorical
                     {:apple (/ 3 4),
                      :orange (/ 1 4)})
                    fruit))
        bin))
```

w = 1.0

fruit = :orange

w = 1.0

:blue

```
(defquery puz1 [fruit]
  (let [bin (sample
              (categorical
               {:red (/ 1 6),
                :blue (/ 5 6)})])
        (if (= bin :red)
            (observe (categorical
                      {:apple (/ 2 8),
                       :orange (/ 6 8)})
                    fruit)
            (observe (categorical
                      {:apple (/ 3 4),
                       :orange (/ 1 4)})
                    fruit))
        bin))
```

```
(defquery puz1 [fruit]
  (let [bin (sample
              (categorical
                {:red (/ 1 6),
                 :blue (/ 5 6)}))]
    (if (= bin :red)
      (observe (categorical
                 {:apple (/ 2 8),
                  :orange (/ 6 8)})
              fruit)
      (observe (categorical
                 {:apple (/ 3 4),
                  :orange (/ 1 4)})
              fruit))
    bin))
```

fruit = :orange

w = 1.0

bin = :blue

fruit = :orange

w = 1.0

bin = :blue

```
(defquery puz1 [fruit]
  (let [bin (sample
              (categorical
               {:red (/ 1 6),
                :blue (/ 5 6)})]]
    (if (= bin :red)
      (observe (categorical
                 {:apple (/ 2 8),
                  :orange (/ 6 8)})
              fruit)
      (observe (categorical
                 {:apple (/ 3 4),
                  :orange (/ 1 4)})
              fruit))
    bin))
```

fruit = :orange

w = 1.0

bin = :blue

```
(defquery puz1 [fruit]
  (let [bin (sample
              (categorical
                {:red (/ 1 6),
                 :blue (/ 5 6)})])
        (if (= bin :red)
            (observe (categorical
                      {:apple (/ 2 8),
                       :orange (/ 6 8)})
                    fruit)
            (observe (categorical
                      {:apple (/ 3 4),
                       :orange (/ 1 4)})
                    fruit))
        bin))
```

fruit = :orange

w = 1.0

bin = :blue

```
(defquery puz1 [fruit]
  (let [bin (sample
              (categorical
               {:red (/ 1 6),
                :blue (/ 5 6)})])
        (if (= bin :red)
            (observe (categorical
                     {:apple (/ 2 8),
                      :orange (/ 6 8)})
                    fruit)
            (observe (categorical
                     {:apple (/ 3 4),
                      :orange (/ 1 4)})
                    fruit))
        bin))
```

```
(defquery puz1 [fruit]
  (let [bin (sample
              (categorical
                {:red (/ 1 6),
                 :blue (/ 5 6)})])
    (if (= bin :red)
      (observe (categorical
                {:apple (/ 2 8),
                 :orange (/ 6 8)})
              fruit)
      (observe (categorical
                {:apple (/ 3 4),
                 :orange (/ 1 4)})
              fruit))
    bin))
```

fruit = :orange

$w = 1.0 \times 0.25$

bin = :blue


```
(defquery puz1 [fruit]
  (let [bin (sample
              (categorical
                {:red (/ 1 6),
                 :blue (/ 5 6)})])
    (if (= bin :red)
      (observe (categorical
                {:apple (/ 2 8),
                 :orange (/ 6 8)})
              fruit)
      (observe (categorical
                {:apple (/ 3 4),
                 :orange (/ 1 4)})
              fruit))
    bin))
```

fruit = :orange

$w = 1.0 \times 0.25$

bin = :blue

Thus, returns
(0.25, :blue)

Likelihood weighted importance sampling

- Simple.
- Regarded as a semi-official semantics for Anglican and other probabilistic PLs.
- Log weight, not weight, used typically. Why?

Likelihood weighted importance sampling

- Simple.
- Regarded as a semi-official semantics for Anglican and other probabilistic PLs.
- Log weight, not weight, used typically. Why?

[Q] OK, but inefficient. Can you guess why?
How can we improve it?

Likelihood weighted importance sampling

[Goal] Estimate $\mathbb{E}_{p(x|y)}[f(x)]$ for a given f .

1. Sample x_1, \dots, x_N from prior $p(x)$.
2. Compute weight $w_i = p(y|x_i)$ for each i .
3. Return weighted avg. $(\sum_i w_i f(x_i)) / \sum_j w_j$.

General ~~likelihood weighted~~

importance sampling

[Goal] Estimate $\mathbb{E}_{p(x|y)}[f(x)]$ for a given f .

1. Sample x_1, \dots, x_N from prior $p(x)$.
2. Compute weight $w_i = p(y|x_i)$ for each i .
3. Return weighted avg. $(\sum_i w_i f(x_i)) / \sum_j w_j$.

General ~~likelihood weighted~~

importance sampling

[Goal] Estimate $\mathbb{E}_{p(x|y)}[f(x)]$ for a given f .

1. Sample x_1, \dots, x_N from ~~prior $p(x)$~~ .
proposal $q(x)$
2. Compute weight $w_i = p(y|x_i)$ for each i .
3. Return weighted avg. $(\sum_i w_i x f(x_i)) / \sum_j w_j$.

General ~~likelihood weighted~~

importance sampling

[Goal] Estimate $\mathbb{E}_{p(x|y)}[f(x)]$ for a given f .

1. Sample x_1, \dots, x_N from ~~prior $p(x)$~~ .
proposal $q(x)$
2. Compute weight $w_i = \frac{p(y|x_i)}{p(y|x_i)p(x_i)/q(x)}$ for each i .
3. Return weighted avg. $(\sum_i w_i x f(x_i)) / \sum_j w_j$.

General ~~likelihood weighted~~

importance sampling

[Goal] Estimate $\mathbb{E}_{p(x|y)}[f(x)]$ for a given f .

1. Sample x_1, \dots, x_N from ~~prior $p(x)$~~ .
proposal $q(x)$
2. Compute weight $w_i = \cancel{p(y|x_i)}$ for each i .
 $p(y|x_i)p(x_i)/q(x)$
3. Return weighted avg. $(\sum_i w_i x f(x_i)) / \sum_j w_j$.

[Q] Why works? Which $q(x)$ is good?

Summary

- Learnt posterior inference using Bayes' rule in the context of discrete probabilities.
- In Anglican, we can condition using observe and perform posterior inference.
- Discussed the likelihood weighted importance sampling algorithm.