

CS492: Probabilistic Programming

# Amortised Inference

Hongseok Yang  
KAIST

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r)))
```

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

Importance sampling with proposal  $q(r)$ .

1. Generate  $(w_1, r_1), \dots, (w_n, r_n)$  by running prog.

2. Estimate  $\mathbb{E}_{p(r|a,b,c)}[f(r)] \approx \sum_i f(r_i) * (w_i / \sum_j w_j)$ .

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

Importance sampling with proposal  $q(r)$ .

1. Generate  $(w_1, r_1), \dots, (w_n, r_n)$  by running prog.

2. Estimate  $\mathbb{E}_{p(r|a,b,c)}[f(r)] \approx \sum_i f(r_i) * (w_i / \sum_j w_j)$ .

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

$w_i = 1$

Importance sampling with proposal  $q(r)$ .

1. Generate  $(w_1, r_1), \dots, (w_n, r_n)$  by running prog.

2. Estimate  $\mathbb{E}_{p(r|a,b,c)}[f(r)] \approx \sum_i f(r_i) * (w_i / \sum_j w_j)$ .



```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

$$w_i = 1$$

Importance sampling with **proposal**  $q(r)$ .

1. Generate  $(w_1, r_1), \dots, (w_n, r_n)$  by running prog.

2. Estimate  $\mathbb{E}_{p(r|a,b,c)}[f(r)] \approx \sum_i f(r_i) * (w_i / \sum_j w_j)$ .

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

$$w_1 = 1 * p(.4)/q(.4)$$
$$r_1 = .4$$

Importance sampling with **proposal**  $q(r)$ .

1. Generate  $(w_1, r_1), \dots, (w_n, r_n)$  by running prog.

2. Estimate  $\mathbb{E}_{p(r|a,b,c)}[f(r)] \approx \sum_i f(r_i) * (w_i / \sum_j w_j)$ .

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

$$w_1 = .096 * p(.4)/q(.4)$$
$$r_1 = .4$$

Importance sampling with proposal  $q(r)$ .

1. Generate  $(w_1, r_1), \dots, (w_n, r_n)$  by running prog.

2. Estimate  $\mathbb{E}_{p(r|a,b,c)}[f(r)] \approx \sum_i f(r_i) * (w_i / \sum_j w_j)$ .

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r)))
```

$$w_1 = .096 * p(.4)/q(.4)$$
$$r_1 = .4$$

Importance sampling with proposal  $q(r)$ .

1. Generate  $(w_1, r_1), \dots, (w_n, r_n)$  by running prog.

2. Estimate  $\mathbb{E}_{p(r|a,b,c)}[f(r)] \approx \sum_i f(r_i) * (w_i / \sum_j w_j)$ .

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

$$w_1 = .096 * p(.4)/q(.4)$$
$$r_1 = .4$$

$$w_2 = .144 * p(.6)/q(.6)$$
$$r_2 = .6$$

Importance sampling with proposal  $q(r)$ .

1. Generate  $(w_1, r_1), \dots, (w_n, r_n)$  by running prog.

2. Estimate  $\mathbb{E}_{p(r|a,b,c)}[f(r)] \approx \sum_i f(r_i) * (w_i / \sum_j w_j)$ .

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

$$w_1 = .096 * p(.4)/q(.4)$$
$$r_1 = .4$$

$$w_2 = .144 * p(.6)/q(.6)$$
$$r_2 = .6$$

Importance sampling with proposal  $q(r)$ .

1. Generate  $(w_1, r_1), \dots, (w_n, r_n)$  by running prog.

2. Estimate  $\mathbb{E}_{p(r|a,b,c)}[f(r)] \approx \sum_i f(r_i) * (w_i / \sum_j w_j)$ .

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

$$w_1 = .096 * p(.4)/q(.4)$$
$$r_1 = .4$$

$$w_2 = .144 * p(.6)/q(.6)$$
$$r_2 = .6$$

Importance sampling with **proposal  $q(r)$** .

1. Generate  $(w_1, r_1), \dots, (w_n, r_n)$  by running prog.

2. Estimate  $\mathbb{E}_{p(r|a,b,c)}[f(r)] \approx \sum_i f(r_i) * (w_i / \sum_j w_j)$ .

**How to find good  $q$ ?**

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (observe (flip r) true)  
        b (observe (flip r) true)  
        c (observe (flip r) false)]  
    r))
```

$$w_1 = .096 * p(.4)/q(.4)$$
$$r_1 = .4$$

$$w_2 = .144 * p(.6)/q(.6)$$
$$r_2 = .6$$

Importance sampling with proposal  $q(r)$ .

1. Generate  $(w_1, r_1), \dots, (w_n, r_n)$  by running prog.

2. Estimate  $\mathbb{E}_{p(r|a,b,c)}[f(r)] \approx \sum_i f(r_i) * (w_i / \sum_j w_j)$ .

How to find good  $q$ ? **Use amortised inference!**



We often need to infer posterior of one model multiple times with different observations.

```
(defquery biased-coin []  
  (let [r (sample (uniform-continuous 0 1))]  
    a (observe (flip r) true)  
    b (observe (flip r) true)  
    c (observe (flip r) false)]  
    r))
```

We often need to infer posterior of one model multiple times with different observations.

```
(defquery biased-coin []  
  (let [r (sample (uniform-continuous 0 1))]  
    a (observe (flip r) false)  
    b (observe (flip r) false)  
    c (observe (flip r) false)]  
    r)))
```

We often need to infer posterior of one model multiple times with different observations.

```
(defquery biased-coin []  
  (let [r (sample (uniform-continuous 0 1))]  
    a (observe (flip r) true)  
    b (observe (flip r) true)  
    c (observe (flip r) true)]  
    r))
```

We often need to infer posterior of one model multiple times with different observations.

```
(defquery biased-coin []  
  (let [r (sample (uniform-continuous 0 1))]  
    a (observe (flip r) false)  
    b (observe (flip r) false)  
    c (observe (flip r) true)]  
    r))
```

We often need to infer posterior of one model multiple times with different observations.

```
(defquery biased-coin []  
  (let [r (sample (uniform-continuous 0 1))  
        a (observe (flip r) false)  
        b (observe (flip r) false)  
        c (observe (flip r) true)]  
    r))
```

Other examples:  
Financial model,  
captcha, brain, etc.

We often need to infer posterior of one model multiple times with different observations.

```
(defquery biased-coin []  
  (let [r (sample (uniform-continuous 0 1))  
        a (observe (flip r) false)  
        b (observe (flip r) false)  
        c (observe (flip r) true)]  
    r))
```

**Other examples:**  
Financial model,  
captcha, brain, etc.

**Amortised inference.**

We often need to infer posterior of one model multiple times with different observations.

```
(defquery biased-coin []  
  (let [r (sample (uniform-continuous 0 1))  
        a (observe (flip r) false)  
        b (observe (flip r) false)  
        c (observe (flip r) true)]  
    r))
```

Other examples:  
Financial model,  
captcha, brain, etc.

Amortised inference. 1) Learn a proposal  $q(x; y)$  parameterized by obs.  $y$  via preprocessing.

We often need to infer posterior of one model multiple times with different observations.

```
(defquery biased-coin []  
  (let [r (sample (uniform-continuous 0 1))  
        a (observe (flip r) false)  
        b (observe (flip r) false)  
        c (observe (flip r) true)]  
    r))
```

Other examples:  
Financial model,  
captcha, brain, etc.

Amortised inference. 1) Learn a proposal  $q(x; y)$  parameterized by obs.  $y$  via preprocessing. 2) Use  $q(x; y_0)$  for any actual observation  $y_0$  later.



We often need to infer posterior of one model multiple times with different observations.

```
(defquery biased-coin []  
  (let [r (sample (uniform-continuous 0 1))  
        a (observe (flip r) false)  
        b (observe (flip r) false)  
        c (observe (flip r) true)]  
    r))
```

Other examples:  
Financial model,  
captcha, brain, etc.

Amortised inference. 1) Learn a proposal  $q(x; y)$  parameterized by obs.  $y$  via preprocessing. 2) Use  $q(x; y_0)$  for any actual observation  $y_0$  later.

neural nets

We often need to infer posterior of one model multiple times with different observations.

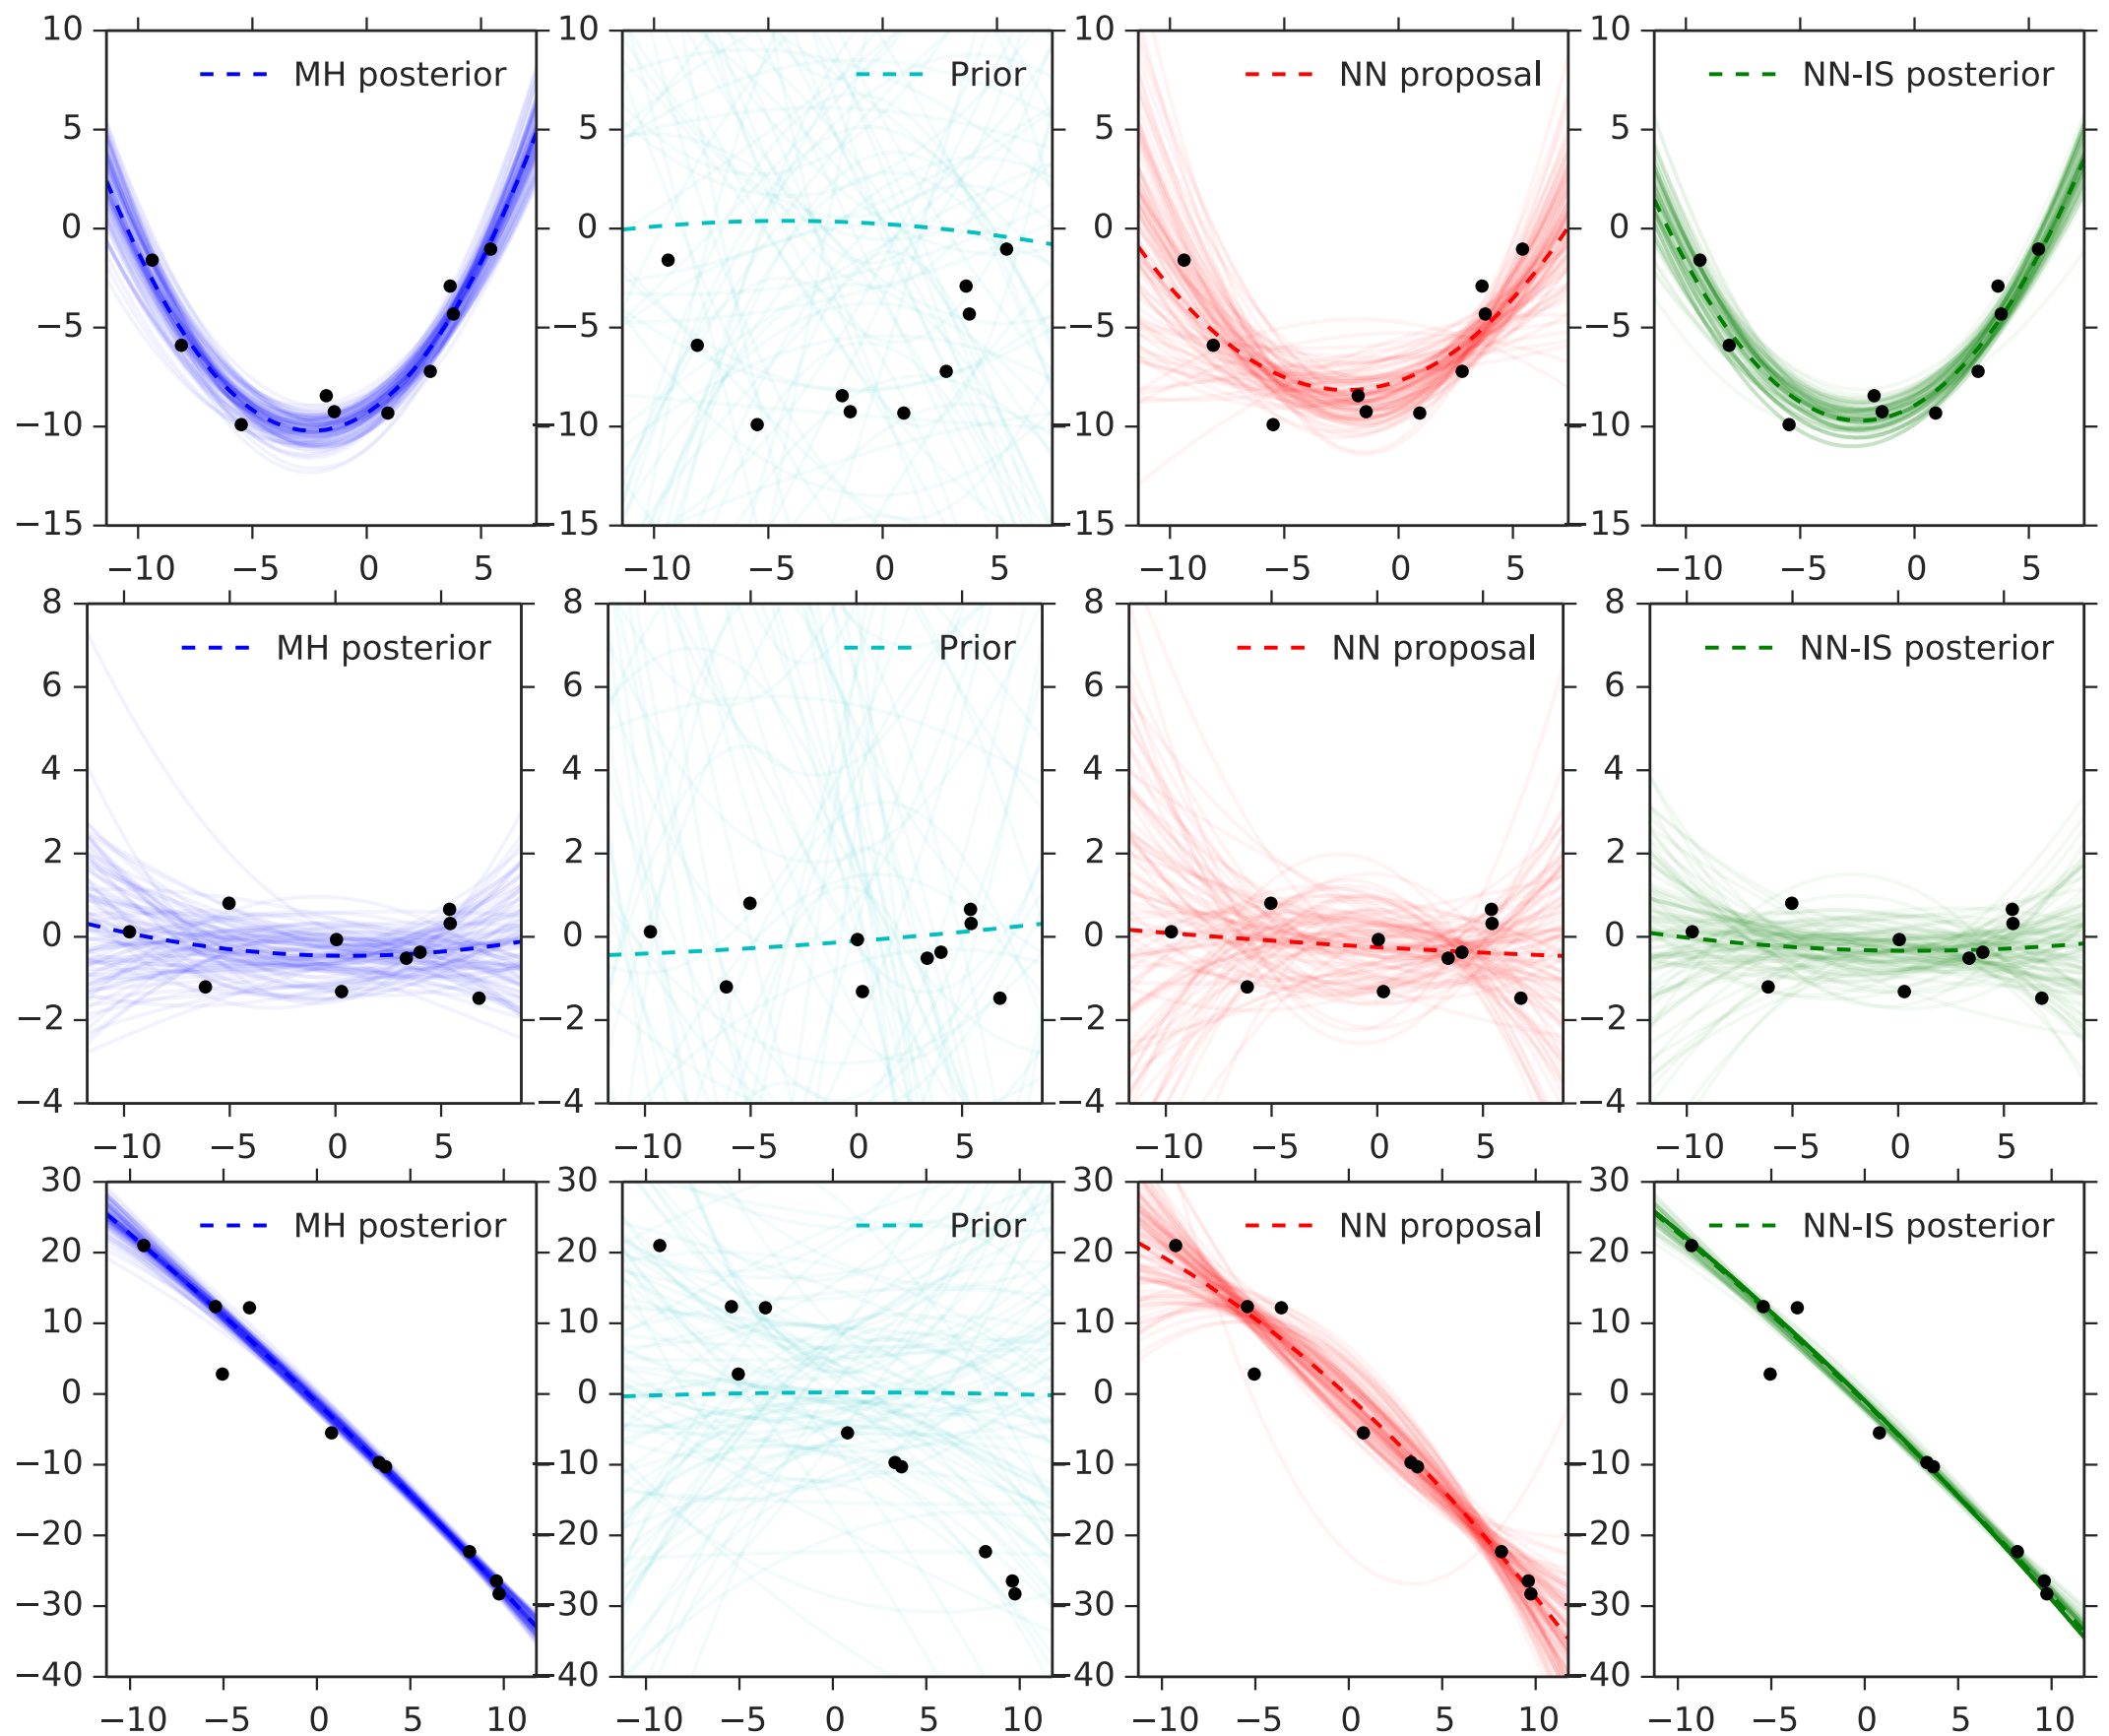
```
(defquery biased-coin []  
  (let [r (sample (uniform-continuous 0 1))  
        a (observe (flip r) false)  
        b (observe (flip r) false)  
        c (observe (flip r) true)]
```

sample/object duality  
and reverse KL

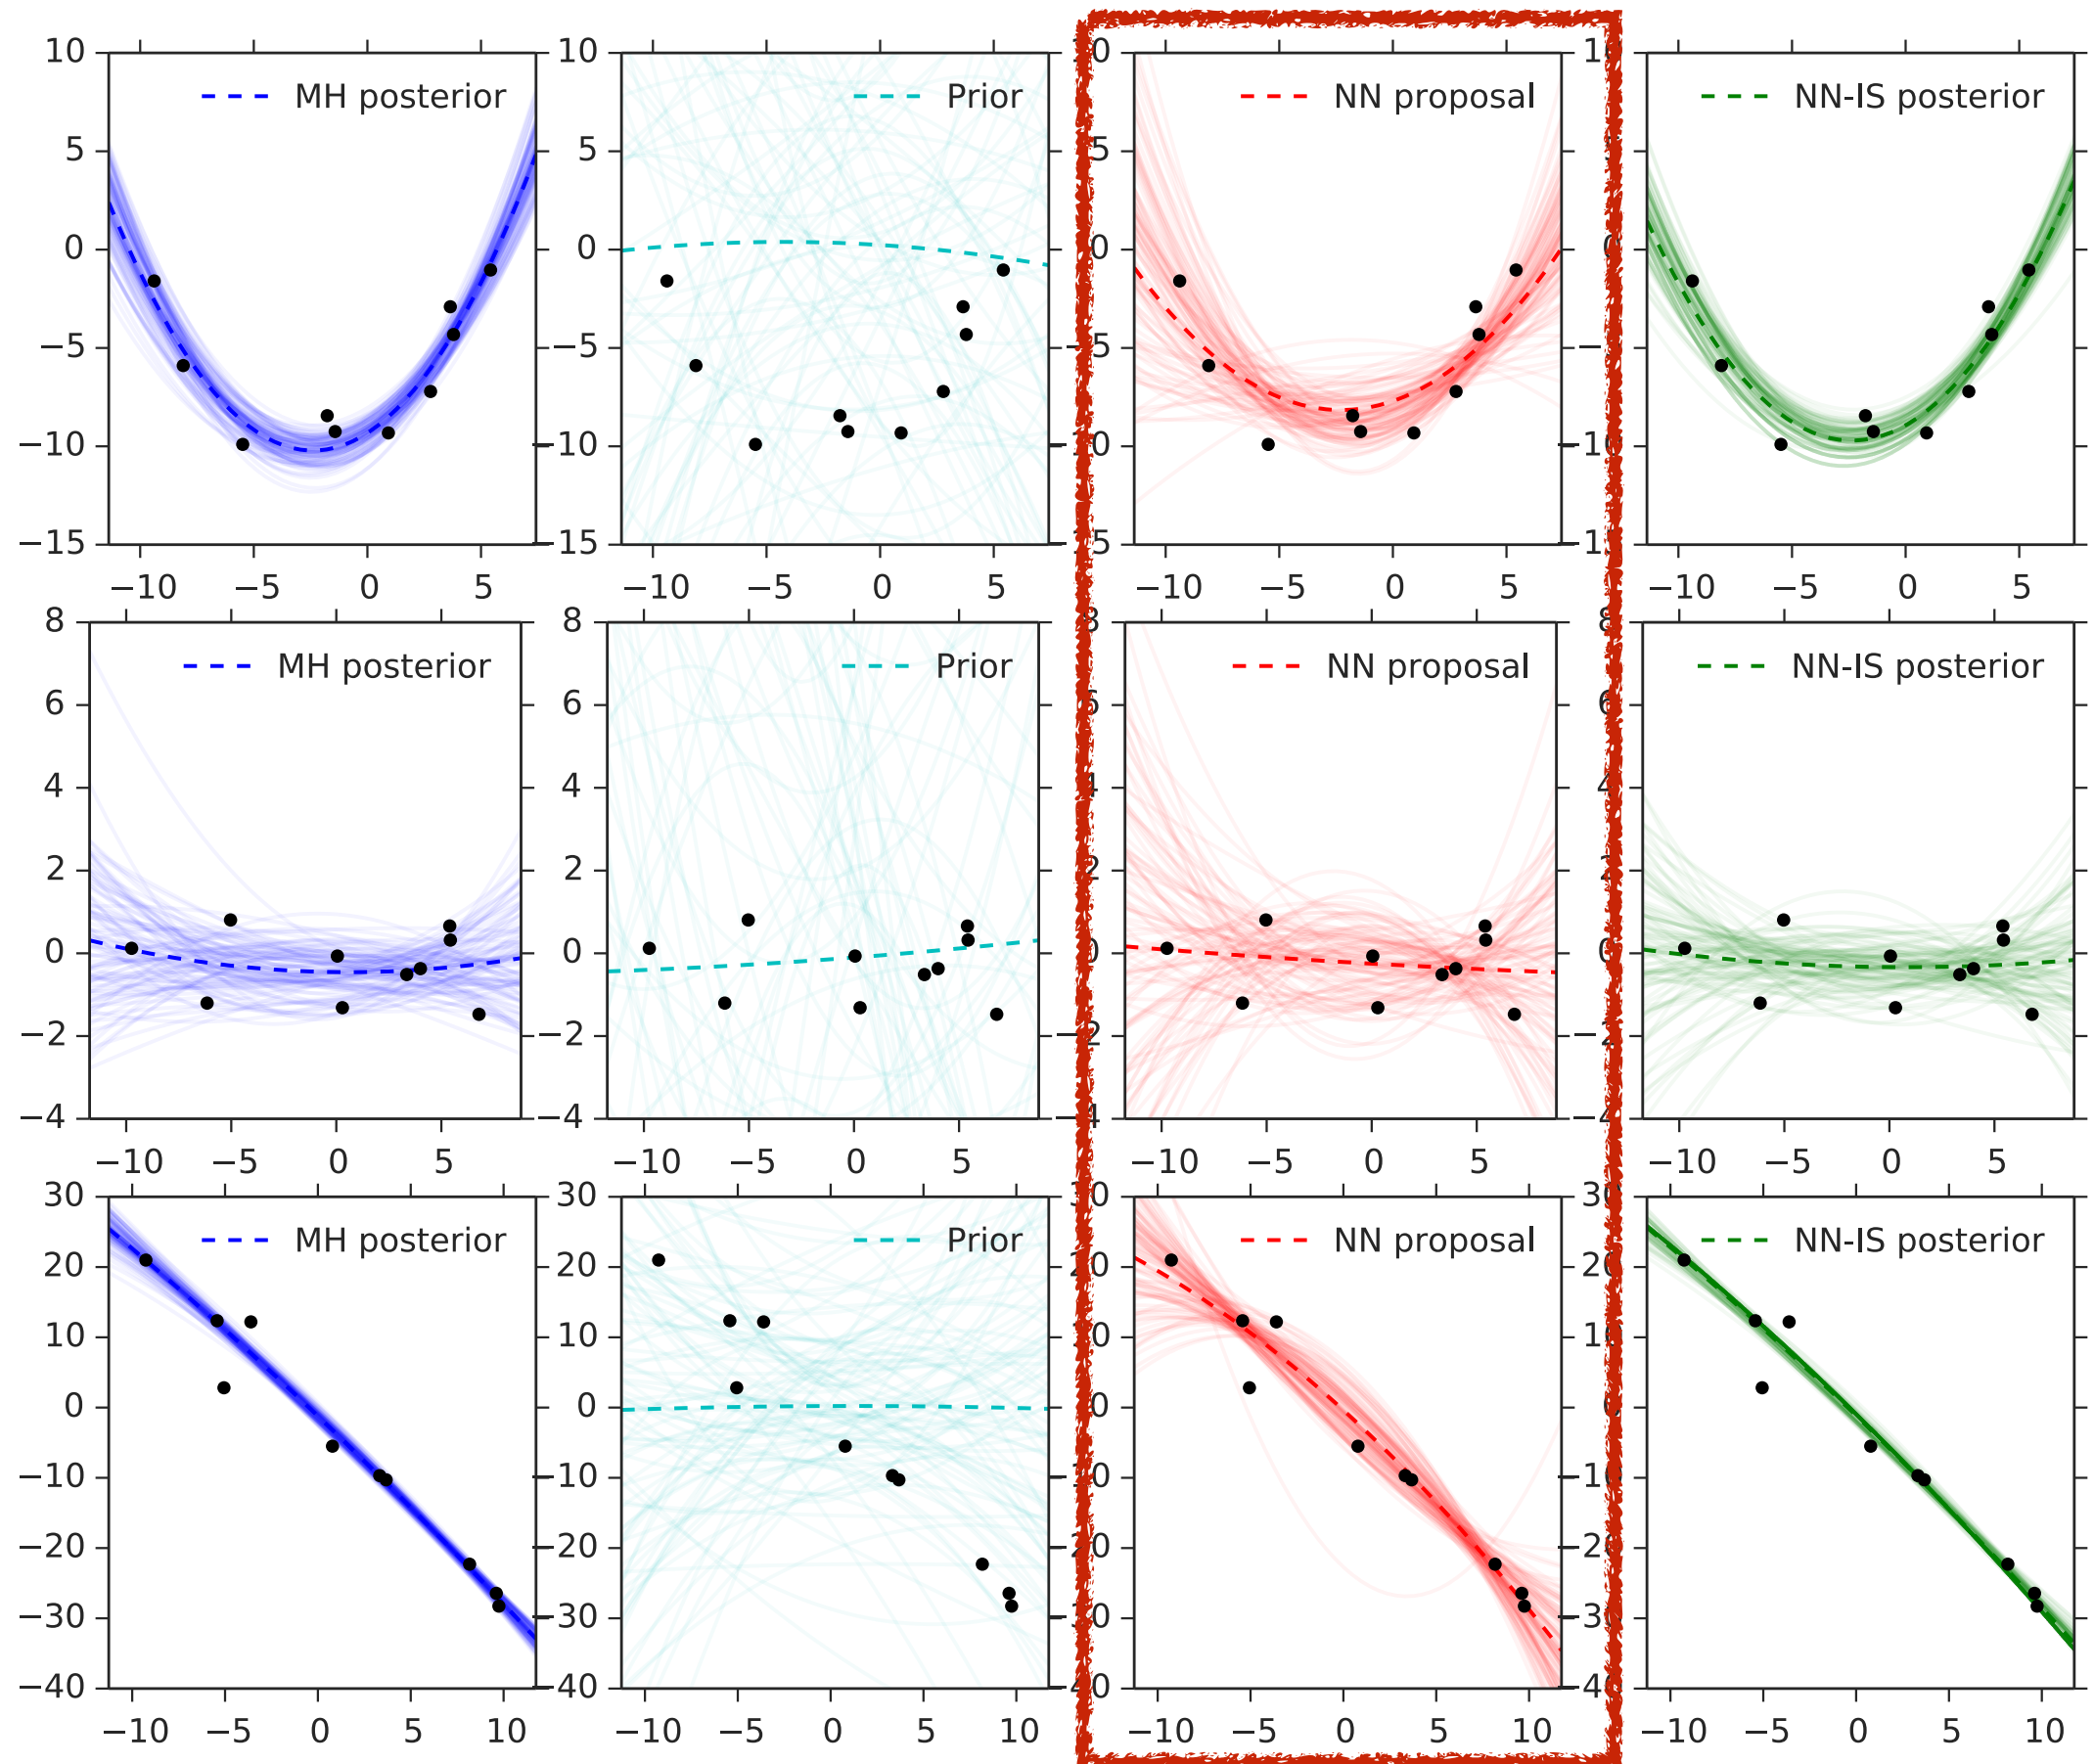
Other examples:  
Financial model,  
captcha, brain, etc.

Amortised inference. 1) **Learn** a proposal  $q(x; y)$  parameterized by obs.  $y$  via preprocessing. 2) Use  $q(x; y_0)$  for any actual observation  $y_0$  later.

neural nets



Model for non-linear regression [Paige et al., ICML16]



Model for non-linear regression [Paige et al., ICML16]



# Observed images

~~W4kgvQ~~  
(W4kgvQ)

~~uV7FeWB~~  
(uV7FeWB)

~~MqhnpT~~  
(MqhnpT)

more preprocessing

# Samples

W4kgvQ

uV7EeWB

MqhnpT

WA4rjvQ

uV7FeWB

MypppT

Woxewd9

mTTEMMm

RIrpES

BKvu2Q

C9QDsoN

rS5FP2B

less preprocessing

Captcha solving [Le et al., AISTATS16]

# Learning outcome

Can describe how amortised inference works for models written in math.

Can explain key ideas behind implementing amortised inference for probabilistic programs.

# Proposal learning problem

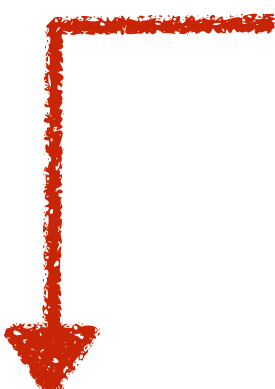
Given:

1. joint dist.  $p(x,y)$  for latent  $x$  and observed  $y$ ,
2. IS proposal  $q_{\theta}(x;y)$  parameterized by  $\theta$  &  $y$ .

Find  $\theta$  such that  $q_{\theta}(x;y)$  is good for most  $y$ .

# Proposal learning problem

Given:



Specified by  $p(x)$  and  $p(y|x)$ .  
Interested in  $p(x|y)$ .  
But specific  $y$  not given yet.

1. joint dist.  $p(x,y)$  for latent  $x$  and observed  $y$ ,
2. IS proposal  $q_{\theta}(x;y)$  parameterized by  $\theta$  &  $y$ .

Find  $\theta$  such that  $q_{\theta}(x;y)$  is good for most  $y$ .



# Proposal learning problem

Given:

1. joint dist.  $p(x,y)$  for latent  $x$  and observed  $y$ ,
2. IS proposal  $q_\theta(x;y)$  parameterized by  $\theta$  &  $y$ .

Find  $\theta$  such that  $q_\theta(x;y)$  is good for most  $y$ .

Differentiable wrt.  $\theta$  for fixed  $x,y$ .

E.g.  $q_\theta(x;y) = \text{normal}(x; f_\theta(y), g_\theta(y))$  for neural nets  $f,g$ .

# Proposal learning problem

Given:

1. joint dist.  $p(x,y)$  for latent  $x$  and observed  $y$ ,
2. IS proposal  $q_{\theta}(x;y)$  parameterized by  $\theta$  &  $y$ .

Find  $\theta$  such that  $q_{\theta}(x;y)$  is good for most  $y$ .

# Proposal learning problem

Given:

1. joint dist.  $p(x,y)$  for latent  $x$  and observed  $y$ ,
2. IS proposal  $q_{\theta}(x;y)$  parameterized by  $\theta$  &  $y$ .

Find  $\theta$  such that  $q_{\theta}(x;y)$  is good for most  $y$ .

# Proposal learning problem

Given:

1. joint dist.  $p(x,y)$  for latent  $x$  and observed  $y$ ,
2. IS proposal  $q_{\theta}(x;y)$  parameterized by  $\theta$  &  $y$ .

Find  $\theta$  such that  $q_{\theta}(x;y)$  is good for most  $y$ .

# Proposal learning problem

Given:

1. joint dist.  $p(x,y)$  for latent  $x$  and observed  $y$ ,
2. IS proposal  $q_{\theta}(x;y)$  parameterized by  $\theta$  &  $y$ .

Find  $\theta$  such that  $q_{\theta}(x;y)$  is good for most  $y$ .

# Proposal learning problem tackled by amortised inf.

Given:

1. joint dist.  $p(x,y)$  for latent  $x$  and observed  $y$ ,
2. IS proposal  $q_{\theta}(x;y)$  parameterized by  $\theta$  &  $y$ .

Find  $\theta$  such that  $q_{\theta}(x;y)$  is good for most  $y$ .

# Proposal learning problem tackled by amortised inf.

Given:

$y$  sampled  
from  $p(y)$

1. joint dist.  $p(x,y)$  for latent  $x$  and observed  $y$ ,
2. IS proposal  $q_{\theta}(x;y)$  parameterized by  $\theta$  &  $y$ .

Find  $\theta$  such that  $q_{\theta}(x;y)$  is good for most  $y$ . ←

# Proposal learning problem tackled by amortised inf.

Given:

$y$  sampled  
from  $p(y)$

1. joint dist.  $p(x,y)$  for latent  $x$  and observed  $y$ ,
2. IS proposal  $q_\theta(x;y)$  parameterized by  $\theta$  &  $y$ .

Find  $\theta$  such that  $q_\theta(x;y)$  is good for most  $y$ .

Small KL divergence from  $p(x|y)$  to  $q_\theta(x;y)$ .

$$\text{KL}[p(x|y) \parallel q_\theta(x;y)] = \mathbb{E}_{p(x|y)}[\log(p(x|y)/q_\theta(x;y))].$$



# Proposal learning problem

$\operatorname{argmin}_{\theta} \mathbb{E}_{p(y)} [\text{KL}[p(x|y) \parallel q_{\theta}(x;y) ]]$ .

Solve this by stochastic gradient descent.

**inf.**

y sampled  
from  $p(y)$

Given:

1. joint dist.  $p(x,y)$  for latent  $x$  and observed  $y$ ,
2. IS proposal  $q_{\theta}(x;y)$  parameterized by  $\theta$  &  $y$ .

Find  $\theta$  such that  $q_{\theta}(x;y)$  is good for most  $y$ .

Small KL divergence from  $p(x|y)$  to  $q_{\theta}(x;y)$ .

$$\text{KL}[p(x|y) \parallel q_{\theta}(x;y)] = \mathbb{E}_{p(x|y)} [\log(p(x|y)/q_{\theta}(x;y))].$$

# Proposal learning problem

$\operatorname{argmin}_{\theta} \mathbb{E}_{p(y)} [\text{KL}[p(x|y) \parallel q_{\theta}(x;y) ]]$ .

Solve this by **stochastic gradient descent**.

inf.

y sampled  
from  $p(y)$

Given:

1. joint dist.  $p(x,y)$  for latent  $x$  and observed  $y$ ,
2. IS proposal  $q_{\theta}(x;y)$  parameterized by  $\theta$  &  $y$ .

Find  $\theta$  such that  $q_{\theta}(x;y)$  is good for most  $y$ .

Small KL divergence from  $p(x|y)$  to  $q_{\theta}(x;y)$ .

$$\text{KL}[p(x|y) \parallel q_{\theta}(x;y)] = \mathbb{E}_{p(x|y)} [\log(p(x|y)/q_{\theta}(x;y))].$$

Stochastic gradient descent  
for  $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Stochastic gradient descent  
for  $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise  $\theta$

Stochastic gradient descent  
for  $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise  $\theta$

$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

# Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise  $\theta$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

# Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise  $\theta$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

# Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise  $\theta$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

...

(until  $\theta$  doesn't change much)



# Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Learning rate

Initialise  $\theta$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

...

(until  $\theta$  doesn't change much)

# Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise  $\theta$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

...



Can't compute, but can approximate.

# Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise  $\theta$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

...



Can't compute, but can approximate.

Sample  $(x_1, y_1), \dots, (x_n, y_n)$  from  $p(x, y)$ .

# Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise  $\theta$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

...



Can't compute, but can approximate.

Sample  $(x_1, y_1), \dots, (x_n, y_n)$  from  $p(x, y)$ .

$$\nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]] \approx -1/n * \sum_{i=1..n} \nabla_{\theta} \log q_{\theta}(x_i; y_i).$$

# Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise  $\theta$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

...

Hard to sample  $x$  from  $p(x|y)$  for given  $y$ ,  
but easy to sample  $(x,y)$  from  $p(x,y)$ .  
Thus, no problem in sampling.

Can't compute, but can approximate.

Sample  $(x_1, y_1), \dots, (x_n, y_n)$  from  $p(x, y)$ .

$$\nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]] \approx -1/n * \sum_{i=1..n} \nabla_{\theta} \log q_{\theta}(x_i; y_i).$$

# Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise  $\theta$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

...



Can't compute, but can approximate.

Sample  $(x_1, y_1), \dots, (x_n, y_n)$  from  $p(x, y)$ .

$$\nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]] \approx -1/n * \sum_{i=1..n} \nabla_{\theta} \log q_{\theta}(x_i; y_i).$$

Exists since  $q_{\theta}(x_i; y_i)$   
is differentiable.



# Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise  $\theta$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

$$\theta \leftarrow \theta - 0.01 * \nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$$

...

**[Q] Prove that this is an unbiased estimator.**

Can't compute, but can approximate.

Sample  $(x_1, y_1), \dots, (x_n, y_n)$  from  $p(x, y)$ .

$$\nabla_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]] \approx -1/n * \sum_{i=1..n} \nabla_{\theta} \log q_{\theta}(x_i; y_i).$$

# Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise  $\theta$

Repeat the following until  $\theta$  doesn't change much:

1. Sample  $(x_1, y_1), \dots, (x_n, y_n)$  from  $p(x, y)$
2.  $G \leftarrow -1/n * \sum_{i=1..n} \nabla_{\theta} \log q_{\theta}(x_i; y_i)$
3.  $\theta \leftarrow \theta - 0.01 * G$



# Learning IS proposal $q_{\theta}(x;y)$ by amortised inference

Using stochastic gradient descent, solve:

$$\operatorname{argmin}_{\theta} \mathbb{E}_{p(y)}[\text{KL}[p(x|y) \parallel q_{\theta}(x;y)]].$$

# Learning IS proposal $q_{\theta}(x;y)$ by amortised inference

Using stochastic gradient descent, solve:

$$\operatorname{argmin}_{\theta} \mathbb{E}_{p(y)} [\text{KL}[p(x|y) \parallel q_{\theta}(x;y)]].$$

[Q] Differences from stochastic variational inf.?

SVI:  $\operatorname{argmin}_{\theta} \text{KL}[q_{\theta}(x) \parallel p(x|y_0)]$  for a given  $y_0$ .

# Learning IS proposal $q_{\theta}(x;y)$ by amortised inference

Using stochastic gradient descent, solve:

$$\operatorname{argmin}_{\theta} \mathbb{E}_{p(y)} [\text{KL}[p(x|y) \parallel q_{\theta}(x;y)]].$$

[Q] Differences from stochastic variational inf.?

SVI:  $\operatorname{argmin}_{\theta} \text{KL}[q_{\theta}(x) \parallel p(x|y_0)]$  for a given  $y_0$ .

(a)  $\text{KL}[\text{true} \parallel \text{approx}]$  vs  $\text{KL}[\text{approx} \parallel \text{true}]$ .

# Learning IS proposal $q_{\theta}(x;y)$ by amortised inference

Using stochastic gradient descent, solve:

$$\operatorname{argmin}_{\theta} \mathbb{E}_{p(y)} [\text{KL}[p(x|y) \parallel q_{\theta}(x;y)]].$$

[Q] Differences from stochastic variational inf.?

SVI:  $\operatorname{argmin}_{\theta} \text{KL}[q_{\theta}(x) \parallel p(x|y_0)]$  for a given  $y_0$ .

(a)  $\text{KL}[\text{true} \parallel \text{approx}]$  vs  $\text{KL}[\text{approx} \parallel \text{true}]$ .

Choice consistent with IS's  
condition on  $q_{\theta}$ 's support.

# Learning IS proposal $q_{\theta}(x;y)$ by amortised inference

Using stochastic gradient descent, solve:

$$\operatorname{argmin}_{\theta} \mathbb{E}_{p(y)} [\text{KL}[p(x|y) \parallel q_{\theta}(x;y)]].$$

[Q] Differences from stochastic variational inf.?

SVI:  $\operatorname{argmin}_{\theta} \text{KL}[q_{\theta}(x) \parallel p(x|y_0)]$  for a given  $y_0$ .

(a)  $\text{KL}[\text{true} \parallel \text{approx}]$  vs  $\text{KL}[\text{approx} \parallel \text{true}]$ .

Choice consistent with IS's  
condition on  $q_{\theta}$ 's support.

Lets us avoid sampling  
from posterior  $p(x|y_0)$ .

# Learning IS proposal $q_{\theta}(x;y)$ by amortised inference

Using stochastic gradient descent, solve:

$$\operatorname{argmin}_{\theta} \mathbb{E}_{p(y)} [\text{KL}[p(x|y) \parallel q_{\theta}(x;y)]].$$

[Q] Differences from stochastic variational inf.?

SVI:  $\operatorname{argmin}_{\theta} \text{KL}[q_{\theta}(x) \parallel p(x|y_0)]$  for a given  $y_0$ .

(a)  $\text{KL}[\text{true} \parallel \text{approx}]$  vs  $\text{KL}[\text{approx} \parallel \text{true}]$ .

# Learning IS proposal $q_{\theta}(x;y)$ by amortised inference

Using stochastic gradient descent, solve:

$$\operatorname{argmin}_{\theta} \mathbb{E}_{p(y)} [\text{KL}[p(x|y) \parallel q_{\theta}(x;y)]].$$

[Q] Differences from stochastic variational inf.?

SVI:  $\operatorname{argmin}_{\theta} \text{KL}[q_{\theta}(x) \parallel p(x|y_0)]$  for a given  $y_0$ .

(a)  $\text{KL}[\text{true} \parallel \text{approx}]$  vs  $\text{KL}[\text{approx} \parallel \text{true}]$ .

(b) Generated  $y$  vs given  $y_0$ .

# Learning IS proposal $q_{\theta}(x;y)$

Lets us avoid sampling  $x$  from posterior  $p(x|y_0)$  for given  $y_0$ . Just need to sample  $(x,y)$  from joint  $p(x,y)$ .

Using stochastic gradient descent, solve:

$$\operatorname{argmin}_{\theta} \mathbb{E}_{p(y)} [\text{KL}[p(x|y) \parallel q_{\theta}(x;y)]].$$

[Q] Differences from stochastic variational inf.?

SVI:  $\operatorname{argmin}_{\theta} \text{KL}[q_{\theta}(x) \parallel p(x|y_0)]$  for a given  $y_0$ .

(a)  $\text{KL}[\text{true} \parallel \text{approx}]$  vs  $\text{KL}[\text{approx} \parallel \text{true}]$ .

(b) Generated  $y$  vs given  $y_0$ .



# Learning IS proposal $q_{\theta}(x;y)$ by amortised inference

Using stochastic gradient descent, solve:

$$\operatorname{argmin}_{\theta} \mathbb{E}_{p(y)} [\text{KL}[p(x|y) \parallel q_{\theta}(x;y)]].$$

[Q] Differences from stochastic variational inf.?

SVI:  $\operatorname{argmin}_{\theta} \text{KL}[q_{\theta}(x) \parallel p(x|y_0)]$  for a given  $y_0$ .

(a)  $\text{KL}[\text{true} \parallel \text{approx}]$  vs  $\text{KL}[\text{approx} \parallel \text{true}]$ .

(b) Generated  $y$  vs given  $y_0$ .

**What about probabilistic programs?**

# Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise  $\theta$

Repeat the following until  $\theta$  doesn't change:

1. Sample  $(x_1, y_1), \dots, (x_n, y_n)$  from  $p(x, y)$
2.  $G \leftarrow -1/n * \sum_{i=1..n} \nabla_{\theta} \log q_{\theta}(x_i; y_i)$
3.  $\theta \leftarrow \theta - 0.01 * G$

# Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise  $\theta$

Repeat the following until  $\theta$  doesn't change:

1. Sample  $(x_1, y_1), \dots, (x_n, y_n)$  from  $p(x, y)$
2.  $G \leftarrow -1/n * \sum_{i=1..n} \nabla_{\theta} \log q_{\theta}(x_i; y_i)$
3.  $\theta \leftarrow \theta - 0.01 * G$

# Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise  $\theta$

Repeat the following until  $\theta$  doesn't change:

1. Sample  $(x_1, y_1), \dots, (x_n, y_n)$  from  $p(x, y)$
2.  $G \leftarrow -1/n * \sum_{i=1..n} \nabla_{\theta} \log q_{\theta}(x_i; y_i)$
3.  $\theta \leftarrow \theta - 0.01 * G$

How to  
sample  $y$ ?

# Sample/observe duality

To sample observations, just replace sample by observe.

```
(defquery biased-coin []  
  (let [r (sample  
            (uniform-continuous 0 1))]  
    a (observe (flip r) true)  
    b (observe (flip r) true)  
    c (observe (flip r) true)]  
    r))
```

# Sample/observe duality

To sample observations, just replace sample by observe.

```
(defquery biased-coin-joint []  
  (let [r (sample  
            (uniform-continuous 0 1))  
        a (sample (flip r))  
        b (sample (flip r))  
        c (sample (flip r))]  
    [r [a b c]]))
```

# Stochastic gradient descent for $\mathbb{E}_{p(y)}[\text{KL}[p(x|y)||q_{\theta}(x;y)]]$

Initialise  $\theta$

Repeat the following until  $\theta$  doesn't change:

1. Sample  $(x_1, y_1), \dots, (x_n, y_n)$  from  $p(x, y)$
2.  $G \leftarrow -1/n * \sum_{i=1..n} \nabla_{\theta} \log q_{\theta}(x_i; y_i)$
3.  $\theta \leftarrow \theta - 0.01 * G$

Compute during execution.  
Similar to the SVI case.  
Just a new rule for sample.



# Last remark

People also use “amortised inference” to mean parameter sharing via neural net in a distr.

Silly example:

1.  $q(x_1, x_2; \theta_1, \theta_2) = q(x_1; \theta_1)q(x_2; \theta_2)$
2.  $q(x_1, x_2; \theta) = q(x_1; f_\theta(1))q(x_2; f_\theta(2))$

where  $f_\theta$  is a neural net.

# References

1. Inference networks for sequential Monte Carlo in graphical models. Paige et al. ICML'16.
2. Inference compilation and universal probabilistic programming. Le et al. AISTATS'17.